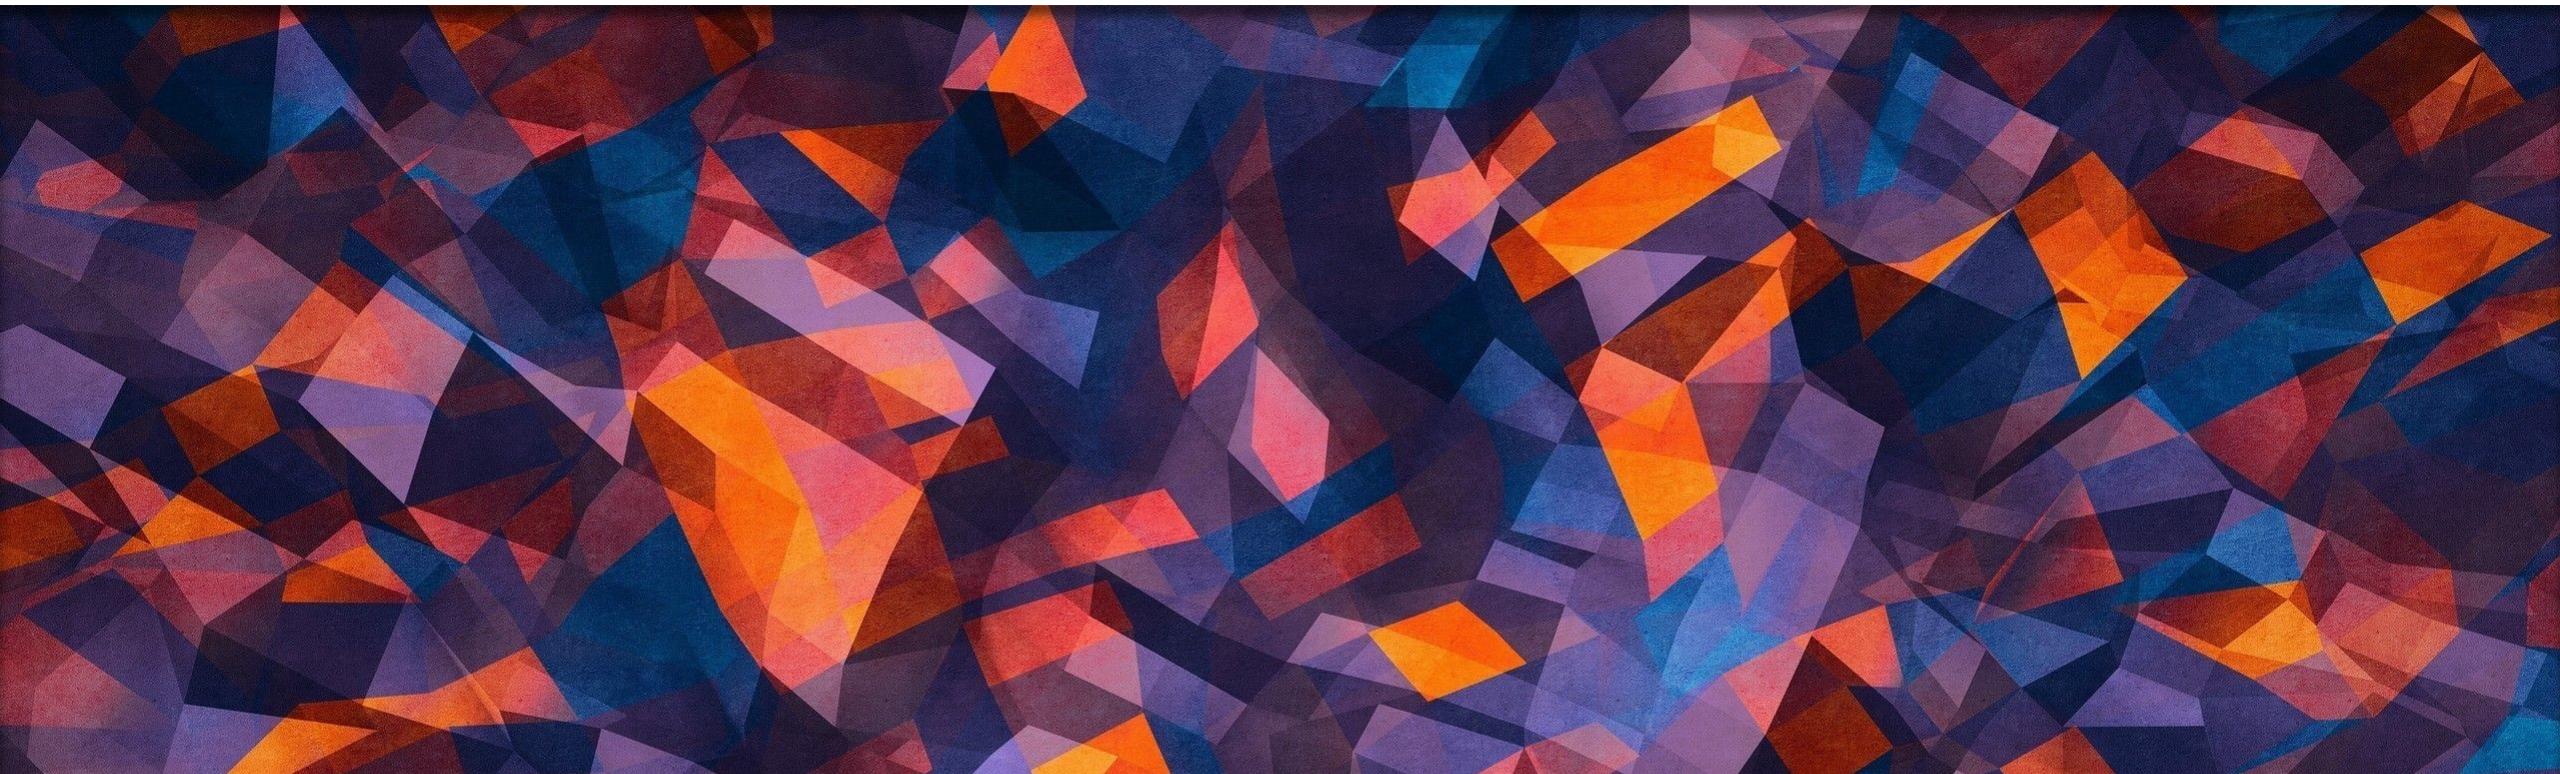




Universität St.Gallen



Web-based Autonomous Systems

# Autonomous Agents: Architectures and Programming Paradigms

Chair for Interaction- and Communication-based Systems (ICS-HSG)

# Our Journey

Prerequisites:  
• ASSE  
• (...)



Week 1:  
**Introduction**

Week 3:  
**Knowledge Representation  
and Reasoning for the Web**

Week 4:  
**Linked Data and Distributed  
Knowledge Graphs**

Week 5:  
**Autonomous Agents  
(Arch. and Programming)**

Week 10:  
**Game Theory and  
Social Choice**

Week 2:  
**A Web for Machines**

Week 7 (Coordination II):  
**Self-Organization and  
Stigmergy**

Exercises

- Ex1: Writing Your First Agent(s)!
- Ex2: Automated Planning
- Ex3: Web Ontologies
- Ex4: Operating on Linked Data
- Ex5: BDI Agents
- Ex6: Interacting Agents on the Web

- Ex7: Ant Colony Optimization
  - Ex8: Organized Agent
  - Ex9: Trustworthy Agents
  - Ex10: Axelrod's Agents
  - Ex11: Reinforcement Learning Agents
- Course Review and Q&A

Week 6 (Coordination I):  
**Agent Communication  
and Interaction**

Week 9 (Coordination IV):  
**Trust & Reputation**

Week 12:  
**An Industry Perspective**



# Our learning objectives so far...

## Learning objectives

- Students understand and can explain the concept of autonomy, how it applies to software systems, and what are the main technological drivers and limitations behind the engineering of autonomous systems.
- Students understand and can explain in depth the design of the Web Architecture, the latest developments, and the tradeoffs and opportunities they bring to the engineering of Web-based autonomous systems.
- Students understand and gain experience with the formal representation and usage of knowledge on the Web.
- Students can design open and large scale Web based autonomous systems - and can explain the implications of their

### Autonomy as a Relational Notion

A multifaceted view drawing from cognitive science [Castelfranchi, 1993; Castelfranchi and Falcone, 2003]

An entity  $X$  is autonomous from  $Y$  about  $P$ , where:

- $X$ : the main entity whose autonomy is considered/evaluated
- $P$ : a function/action/goal that must be realized or maintained by the main entity and on which the autonomy is evaluated
- $Y$ : the secondary entity (human, artificial agent, environment, organization, etc.) with respect to whom  $X$  should be considered autonomous given the specified function/action/goal  $P$

Concrete examples:

- autonomy **from other agents**
- autonomy **from organizations**
- autonomy **from the environment**

C. Castelfranchi. Guarantees for Autonomy in Cognitive Agent Architectures. In: C. Castelfranchi, R. Falcone. *From Automaticity to Autonomy: The Frontier of Artificial Agents*. In: Hexmoor H., Castelfranchi C., Falcone R. (eds) *Autonomy, Multiagent Systems, Artificial Societies, and Simulated Organizations* (International Book Series)

abilities to prototype autonomous systems on the Web and off

With **limitations & state-of-the-art tooling**

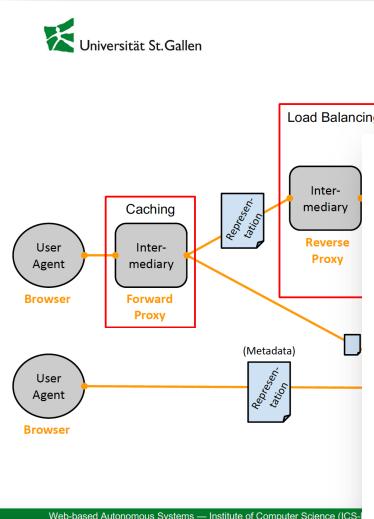
### Technological drivers discussed so far:

- automated planning (“AI planning”)
- knowledge representation and reasoning (knowledge graphs)
- Linked Data
- Re-decentralizing the Web (Solid)

# Our learning objectives so far...

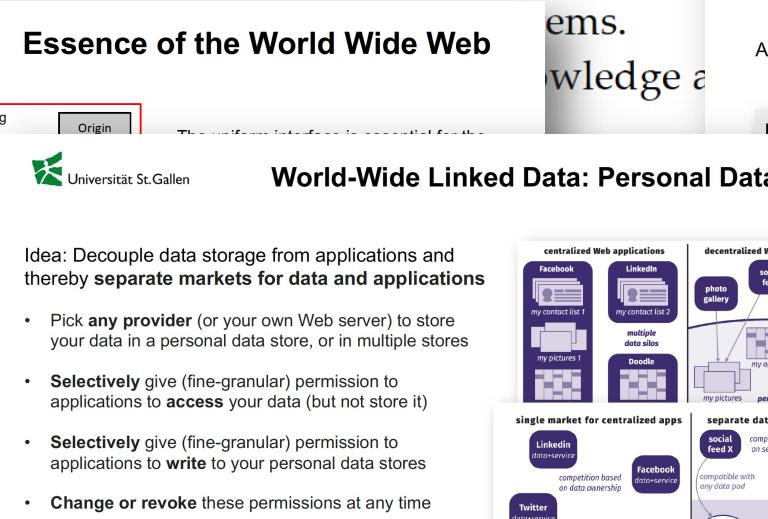
## Learning objectives

- Students understand and can explain the concept of autonomy, how it applies to software systems, and what are the main technological drivers and limitations behind the engineering of autonomous systems.
- Students understand and can explain in depth the design of the Web Architecture, the latest developments, and the tradeoffs and opportunities they bring to the engineering of Web-based autonomous systems.
- Students understand and gain experience with the formal representation and usage of knowledge on the Web.
- Students can design open and large-scale Web-based a



**Essence of the World Wide Web**

The diagram illustrates the flow of data from two User Agents (Browsers) through an Intermediary (containing Caching and Forward Proxy) to a Reverse Proxy, which then connects to an Origin (Load Balancing). The Reverse Proxy also interacts with a Representation layer and a Metadata layer.

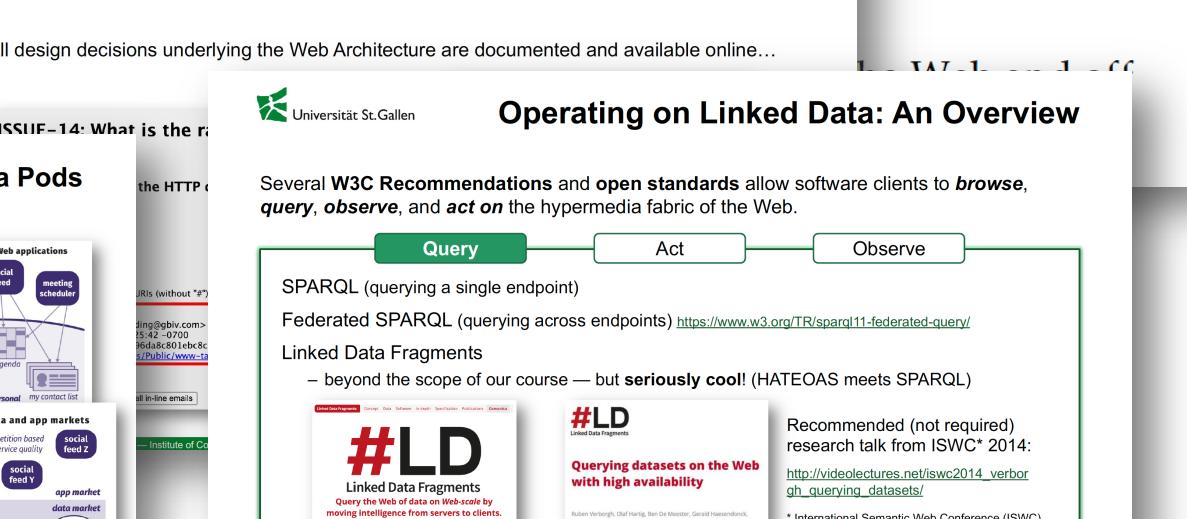


**World-Wide Linked Data: Personal Data Pods**

Idea: Decouple data storage from applications and thereby **separate markets for data and applications**

- Pick **any provider** (or your own Web server) to store your data in a personal data store, or in multiple stores
- **Selectively** give (fine-granular) permission to applications to **access** your data (but not store it)
- **Selectively** give (fine-granular) permission to applications to **write** to your personal data stores
- **Change or revoke** these permissions at any time

**Solid** (Social Linked Data)  
<https://solidproject.org/>



**Operating on Linked Data: An Overview**

Several W3C Recommendations and open standards allow software clients to **browse**, **query**, **observe**, and **act on** the hypermedia fabric of the Web.

**Query**: SPARQL (querying a single endpoint), Federated SPARQL (querying across endpoints) <https://www.w3.org/TR/sparql11-federated-query/>

**Act**: Linked Data Fragments

- beyond the scope of our course — but **seriously cool!** (HATEOAS meets SPARQL)

**Observe**: #LD (Linked Data Fragments), Querying datasets on the Web with high availability

Recommended (not required) research talk from ISWC\* 2014:  
[http://videolectures.net/iswc2014\\_verborgh\\_querying\\_datasets/](http://videolectures.net/iswc2014_verborgh_querying_datasets/)

\* International Semantic Web Conference (ISWC)

# Our learning objectives so far...

## Learning objectives

## URLs, etc

"Centralized resources" on a decentralized web"

- Students understand and can explain the concept of autonomy, how it applies to software systems, and what are the main technological drivers and limitations behind the engineering of autonomous systems.
  - Students understand and can explain in depth the design of the Web Architecture, the latest developments, and the tradeoffs and opportunities they bring to the engineering of Web-based autonomous systems.
  - Students understand and gain experience with the formal representation and usage of knowledge on the Web.
  - Students can design open and large-scale Web-based autonomous systems - and can explain the implications of their

**Ontologies**

An ontology is a formal knowledge representation language [Studer et al., 2009].

using a formal representation language like the Web Ontology Language (OWL).

**Knowledge Graphs**

"A knowledge graph is a **graph of data** intended to accumulate and convey **knowledge of the real world**, whose **nodes** represent **entities of interest** and whose **edges** represent **relations between these entities**" [Hogan et al., 2021].

**Syntax**  
graph of data  
nodes  
edges

**Semantics**  
knowledge entities  
relations between them

**Graph Datasets and Querying**

Graph query languages (SPARQL, Cypher, etc.) share 3 types of **graph patterns**:

- basic graph patterns
- complex graph patterns (relational algebra operators)
- navigational graph patterns (path expressions)

Q: "What are food festivals held in a city reachable from Arica by bus or flight?"

Food Festival → type → ?event → name → ?name  
 Arica → (bus | flight)\* → ?city → inverse  
 or  
 Kleene star

Existential Language  
 Concept intersection ( $C \sqcap D$ ), existential restrictions ( $\exists R.C$ )  
 Attributive Language with Complement  
 Concept intersection ( $C \sqcap D$ ), concept union ( $C \sqcup D$ ), concept negation ( $\neg C$ ), existential restrictions ( $\exists R.C$ )  
 extends ALC with transitive closure (Trans( $R$ ))

**DLs: A Family of Logics**

**Semantic Web: "Consistent Logical Web of Data"**

Triple

RDF

Circles and documents of the same color indicate correspondence between non-information and information resources

Circles of the same color indicate identical resources

Tobias Käfer, Andreas Harth, and Lars Heling, Knowledge Graphs and Linked Data, Al4Industry Summer School, MINES Saint-Étienne, 2021.

# Our learning objectives so far...

## Learning objectives

- Students understand and can explain the concept of autonomy, how it applies to software systems, and what are the main technological drivers and limitations behind the engineering of autonomous systems.
- Students understand and can explain in depth the design of the Web Architecture, the latest developments, and the tradeoffs and opportunities they bring to the engineering of Web-based autonomous systems.
- Students understand and gain experience with the formal representation and usage of knowledge on the Web.
- Students can design open and large-scale Web-based autonomous systems - and can explain the implications of their design choices for the engineering of such systems.
- Students possess the relevant engineering knowledge and abilities to prototype autonomous systems on the Web and off the Web.

*~ broader than interoperability*

**Openness:** The ability of a system to accommodate new components, which may or may not have been foreseen at design time.

**Scalability:** The ability of a system to meet new requirements of size and scope.

*~ broader than accommodate many conc. users*

> TDR (Ontologies, Semantics)  
> Coordination: Join multi-agt sys.,  
how to interact, trust  
> Regulation

Let's discuss!

*SPARQL*  
> Link-traversal: Architecture of the web!  
↳ design multi-agt. sys on web inherit prop. ~

# Our learning objectives so far...

## Learning objectives

- Students understand and can explain the concept of autonomy, how it applies to software systems, and what are the main technological drivers and limitations behind the engineering of autonomous systems.
- Students understand and can explain in depth the design of the Web Architecture, the latest developments, and the tradeoffs and opportunities they bring to the engineering of Web-based autonomous systems.
- Students understand and gain experience with the formal representation and usage of knowledge on the Web.
- Students can design open and large-scale Web-based autonomous systems - and can explain the implications of their design choices for the engineering of such systems.
- Students possess the relevant engineering knowledge and abilities to prototype autonomous systems on the Web and off the Web.

So far, focus on **specific methods** for enabling autonomy

- knowledge representation, reasoning, automated planning, etc.

Today, we'll start prototyping autonomous agents

# Our Journey

Prerequisites:  
• ASSE  
• (...)



Week 1:  
**Introduction**



Week 3:  
**Knowledge Representation  
and Reasoning for the Web**



Week 4:  
**Linked Data and Distributed  
Knowledge Graphs**



Week 5:  
**Autonomous Agents  
(Arch. and Programming)**



Week 10:  
**Game Theory and  
Social Choice**



Week 11:  
**Reinforcement Learning  
and Multi-Agent Learning**



Week 9 (Coordination IV):  
**Trust & Reputation**



Week 6 (Coordination I):  
**Agent Communication  
and Interaction**



Week 2:  
**A Web for Machines**



Week 7 (Coordination II):  
**Self-Organization and  
Stigmergy**



Week 12:  
**An Industry Perspective**



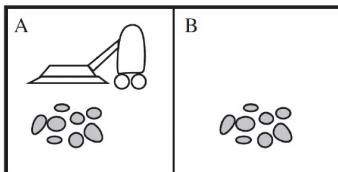
Exercises

Ex1: Writing Your First Agent(s)!  
Ex2: Automated Planning  
Ex3: Web Ontologies  
Ex4: Operating on Linked Data  
Ex5: BDI Agents  
Ex6: Interacting Agents on the Web

Ex7: Ant Colony Optimization  
Ex8: Organized Agent  
Ex9: Trustworthy Agents  
Ex10: Axelrod's Agents  
Ex11: Reinforcement Learning Agents  
Course Review and Q&A

# Lecture #1: The Vacuum-Cleaner Agent Revisited

## Example: A Vacuum-Cleaner Agent



Vacuum-cleaner world with **2 locations** (A and B)

**Percepts:** [location, status], e.g. [A, Clean]

**Actions:** MoveRight, MoveLeft, Vacuum, DoNothing

Partial tabulation of the **agent** function:

## Example: A Vacuum-Cleaner Agent

The **agent function** is implemented by an **agent program** running on an **agent architecture**

abstract mathematical description  
of the agent's behavior

concrete implementation  
of the agent's behavior

more of that Lecture #5!

```
1  function Reflex-Vacuum-Agent([location, status]) returns an action
2    if status == Dirty then return Vacuum
3    else if location == A then return MoveRight
4    else if location == B then return MoveLeft
```

Agent program for a simple vacuum-cleaner agent that implements our agent function.

Reactive  
Agent

We can say a lot about an agent by **observing** its behavior (without looking into the agent's internals)  
– in practice, this is often necessary (different agent architectures, programming languages, etc.)

# Types of Agents (for today's lecture)

Automated Planning

Situation → Action

Action is **coupled** to perception

- agents react to percepts (**reactive behavior**)
- agents are usually programmed in terms of **condition-action rules** "if  $X$  then  $Y$ "
- agents **may or may not** have a model of their environment
- agent communication is usually **indirect** through the environment (stigmergy)

**Reactive Agents**

Situation → Deliberation & Means-End → Action

Action is **decoupled** from perception through deliberation & means-end reasoning

- agents react to percepts and pursue long-term goals (**reactive and proactive behavior**)
- agents are usually programmed in terms of **mental states** (in the **agent-oriented programming** paradigm)
- agents usually **have a model** of their environment
- agent communication is both **direct** through messages and **indirect** through the environment

**Deliberative Agents**

≈ Cognitive Agt.

current state of e.  
goal state of e.

but w/o mental states

Lecture #11:  
reinforcement  
learning agents

Both types of agents are useful depending on the design requirements at hand

In this course, we are taking a software engineering viewpoint (not an AI viewpoint)

# From Software Architecture...

ASSE Review



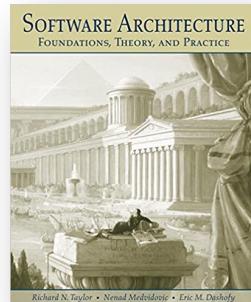
## Software Architecture

**Software architecture** consists of the *structure* of the system, combined with *architecture characteristics* (“-ilities”) the system must support, *architecture decisions*, and finally *design principles*.

Source: Fundamentals of Software Architecture

REST design decisions

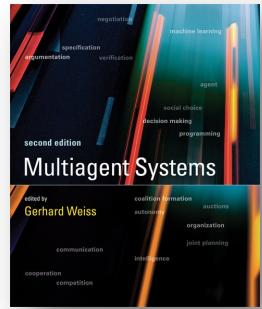
↪  
properties



“A software system’s architecture is **the set of principal design decisions** made about the system.”

— Richard N. Taylor et al., 2010: Software Architecture: Foundations, Theory, and Practice

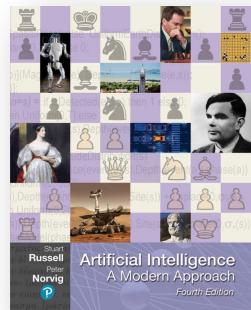
# ...to Agent Architecture



“Agent architectures are **software architectures for decision-making systems** that are **embedded in an environment**.”

— Michael Wooldridge, 2013: Intelligent Agents. In Multiagent Systems (Second Edition).

deliberate, reason  
perceive & act



“The job of AI is to design an agent program that implements the agent function—the mapping from percepts to actions. We assume this program will run on **some sort of computing device with physical sensors and actuators**—we call this the **architecture**: *agent = architecture + program*.”

—Russell and Norvig, 2020: Artificial Intelligence: A Modern Approach (Fourth Edition).

In **software agents**, we think of the **agent architecture** as a **Process Virtual Machine**

create/kill  
ξ  
Process VM: Agt. is process on a VM

# Today's Agenda

- Introduction
- Reactive Agents
  - The Subsumption Architecture
- Deliberative Agents
  - Agent-Oriented Programming
  - The Belief-Desire-Intention (BDI) Architecture
  - Programming BDI Agents in Jason

# Reactive Web (of Things) Systems

## Rule-Based Programming

**IFTTT**

<http://ifttt.com>

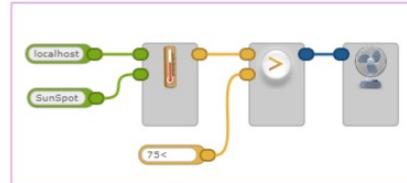
**zapier**

<http://zapier.com>

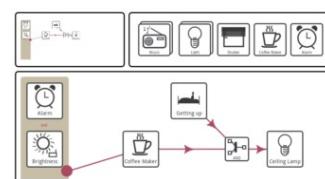
If New feed item from Interactions-HSG website, then Post a tweet to @InteractionsHSG  
by hsginteractions  
Connected  
1  
Twitter RSS

We can use rule-based programming for  
Linked Data as well  
(e.g., see [Käfer & Harth, 2018])

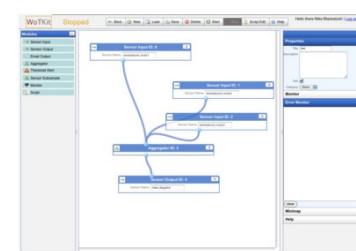
## Dataflow programming for the Web of Things



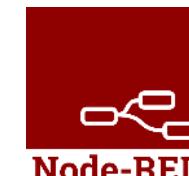
[Guinard et al., 2009]



[Rietzler et al., 2013]



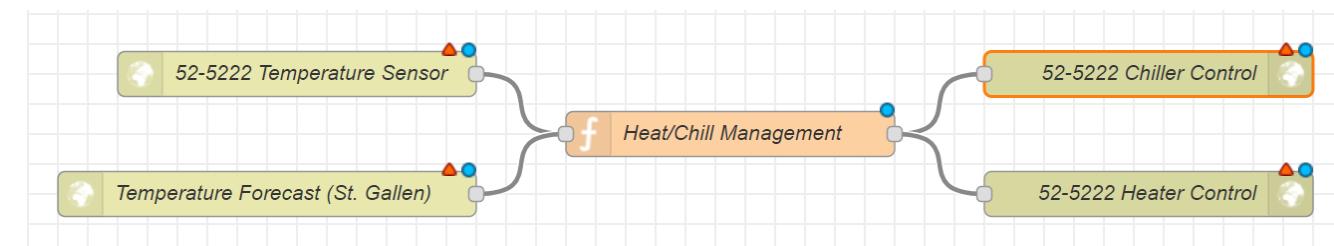
[Blackstock & Lea, 2012]



**Node-RED**

"Node-RED is a programming tool for wiring together hardware devices, APIs and online services in new and interesting ways."

<http://nodered.org>



T. Käfer and A. Harth. Rule-based Programming of User Agents for Linked Data. LDOW 2018.

D. Guinard, V. Trifa, T. Pham, and O. Liechti. Towards physical mashups in the web of things. INSS 2009.

M. Blackstock and R. Lea. 2012. WoTKit: A lightweight toolkit for the web of Things. WoT 2012.

M Rietzler, J. Greim, M. Walch, F. Schaub, B. Wiedersheim, and M. Weber. HomeBLOX: Introducing process-driven home automation. HomeSys 2013.

# Reactive Web (of Things) Systems

Rule-Based Programming

**IFTTT**

<http://ifttt.com>

**zapier**

<http://zapier.com>

If New feed item  
from Interactions-  
HSG website, then  
Post a tweet to  
@InteractionsHSG

by hsginteractions  
Connected  
1  
Twitter Wi-Fi



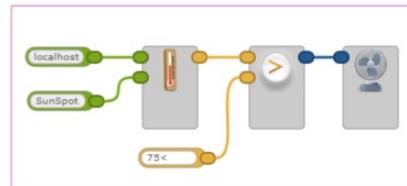
Lecture #1: Ericsson's  
“Social Web of Things”

<https://youtu.be/i5AuzQXBsG4>

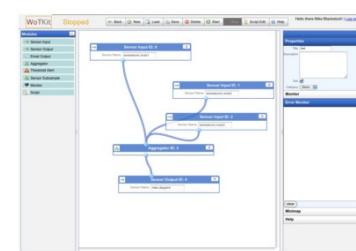
What was the problem?

> Heterogenous systems, complexity

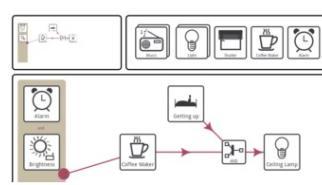
Dataflow programming for the Web of Things



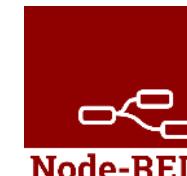
[Guinard et al., 2009]



[Blackstock & Lea, 2012]



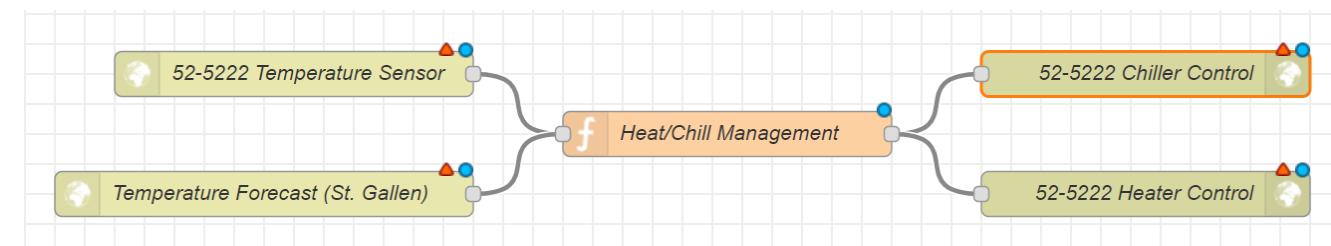
[Rietzler et al., 2013]



Node-RED

“Node-RED is a programming tool for wiring together hardware devices, APIs and online services in new and interesting ways.”

<http://nodered.org>



T. Käfer and A. Harth. Rule-based Programming of User Agents for Linked Data. LDOW 2018.

D. Guinard, V. Trifa, T. Pham, and O. Liechti. Towards physical mashups in the web of things. INSS 2009.

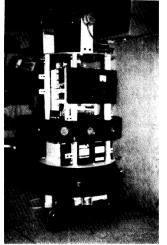
M. Blackstock and R. Lea. 2012. WoTKit: A lightweight toolkit for the web of Things. WoT 2012.

M. Rietzler, J. Greim, M. Walch, F. Schaub, B. Wiedersheim, and M. Weber. HomeBLOX: Introducing process-driven home automation. HomeSys 2013.

# The Subsumption Architecture

**Situated Agents**

Intelligent, rational behavior is **innately linked to the environment** an agent occupies [Wooldridge, 2013]  
 ⇒ important shift in the development of artificially intelligent agents (mid-80s to present)

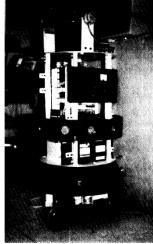


"The MIT AI Lab mobile robot"  
[Brooks, 1985]

Web-based Autonomous Systems — Institute of Computer Science (ICS-HSG)

**Situated Agents**

Intelligent, rational behavior is **innately linked to the environment** an agent occupies [Wooldridge, 2013]  
 ⇒ important shift in the development of artificially intelligent agents (mid-80s to present)



"The MIT AI Lab mobile robot"  
[Brooks, 1985]

Traditional Artificial Intelligence has adopted a style of research where the agents that are built to test theories in intelligence are essentially problem solvers that work in an symbolic abstracted domain. The symbols may have referents in the minds of the builders of the systems, but there is nothing to ground those referents in any real world. Furthermore, the agents are not situated in a world at all. Rather they are given a problem, and they solve it. Then, they are given another problem and they solve it. They are not participating in a world as would agents in the usual sense.

[Brooks, 1991]

Michael Wooldridge. *Intelligent Agents*. In *Multagent Systems*, Second Edition (ed. Gerhard Weiss), 2013.  
 Rodney A. Brooks. *A Robust Layered Control System for a Mobile Robot*. MIT A.I. Laboratory, A.I. Memo No. 864, 1985.  
 Rodney A. Brooks. *Intelligence Without Reason*. MIT A.I. Laboratory, A.I. Memo No. 1293, 1991.

## Lecture #1

- > Back in the days robots took more time to plan → env. chg. changed
- > situatedness: Agt. needs to be in env. And needs to react shortly



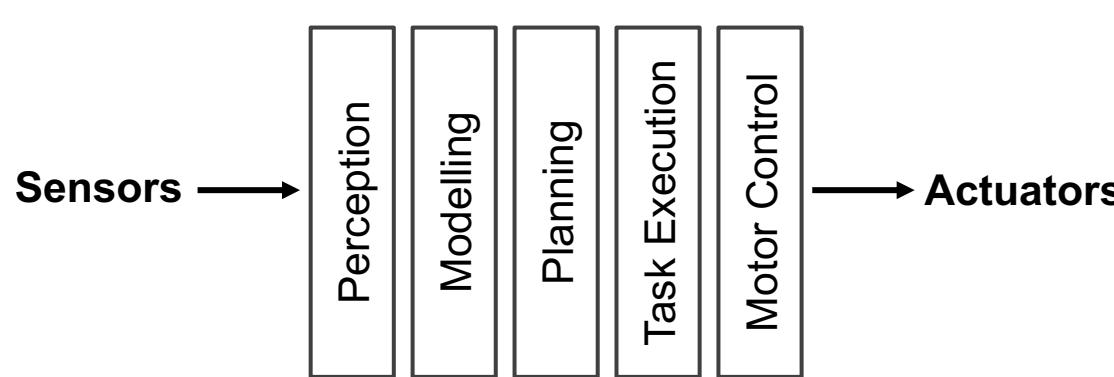
**Rodney Brooks, “Robots Rising”**  
 Discover Channel, 1998 (from 16:43 to 17:37)  
<https://youtu.be/7-Na7nkAH1k?t=1003>

Rodney A. Brooks. *A Robust Layered Control System for a Mobile Robot*. MIT A.I. Laboratory, A.I. Memo No. 864, 1985.  
 Rodney A. Brooks. *Intelligence Without Reason*. MIT A.I. Laboratory, A.I. Memo No. 1293, 1991.

# The Subsumption Architecture

Radical new idea: **concurrent behaviors**

- **levels of competence:** each behavioral level runs independently, concurrently, and asynchronously  
⇒ each behavioral level continually takes perceptual input and maps it to an action (can fire simultaneously)



Decomposition of a mobile robot control system into **functional modules** (horizontal decomposition)

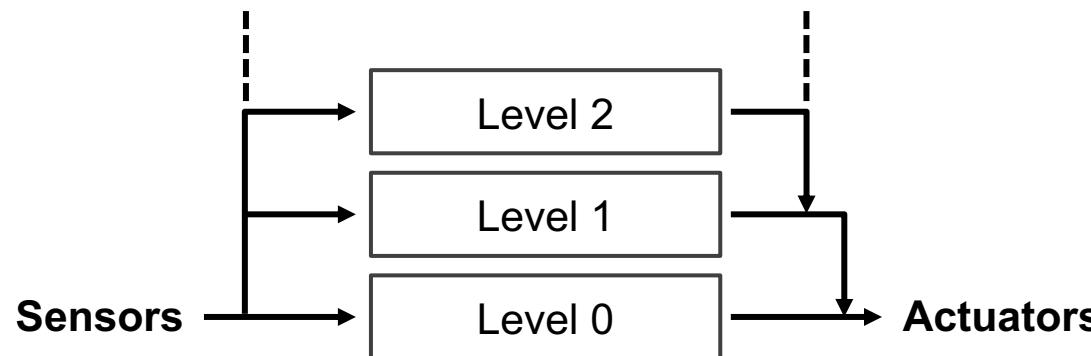


Decomposition of a mobile robot control system into **task achieving behaviors** (vertical decomposition)

# The Subsumption Architecture

## Radical new idea: concurrent behaviors

- **levels of competence:** each behavioral level runs independently, concurrently, and asynchronously  
⇒ each behavioral level continually takes perceptual input and maps it to an action (can fire simultaneously)
- **layers of control:** higher behavioral levels subsume the roles of lower levels in order to take control
  - the system can be partitioned at any level and the lower layers form a complete operational control system



**Control is layered:** higher behavioral levels can **suppress** the input and **inhibit** the output of lower levels



Decomposition of a mobile robot control system into **task achieving behaviors** (vertical decomposition)

# The Subsumption Architecture



(founded in 1990 by Rodney Brooks  
and two MIT graduate students)



A Can-Collecting Robot [Connell, 1989]:  
<https://youtu.be/YtNKuwiVYm0>

- No built-in model of the environment
- No symbolic representation of knowledge
- Behavioral modules as finite state machines

Remember **situation calculus**? (Lecture #2)

Proposed in **strong reaction** against the  
“traditional” logic-based agents

**Highly influential** in robotics, but also for  
the engineering of software agents

Remember **situatedness, embodiment**? (Lecture #1)

Jonathan Connell. *A Colony Architecture for an Artificial Creature*. MIT A.I. Laboratory, Technical Report No. 1151, 1989.

# Today's Agenda

- Introduction
- Reactive Agents
  - The Subsumption Architecture
- Deliberative Agents
  - Agent-Oriented Programming
  - The Belief-Desire-Intention (BDI) Architecture
  - Programming BDI Agents in Jason

## Agent-oriented programming

Yoav Shoham

*Robotics Laboratory, Computer Science Department, Stanford University, Stanford, CA 94305, USA*

Received June 1991

Revised February 1992

### *Abstract*

Shoham, Y., Agent-oriented programming, Artificial Intelligence 60 (1993) 51–92.

A new computational framework is presented, called *agent-oriented programming* (AOP), which can be viewed as a specialization of *object-oriented programming*. The state of an agent consists of components such as beliefs, decisions, capabilities, and obligations; for

### Programming Paradigms & Metaphors

imperative => machines

functional => mathematics

object-oriented => “world of objects”

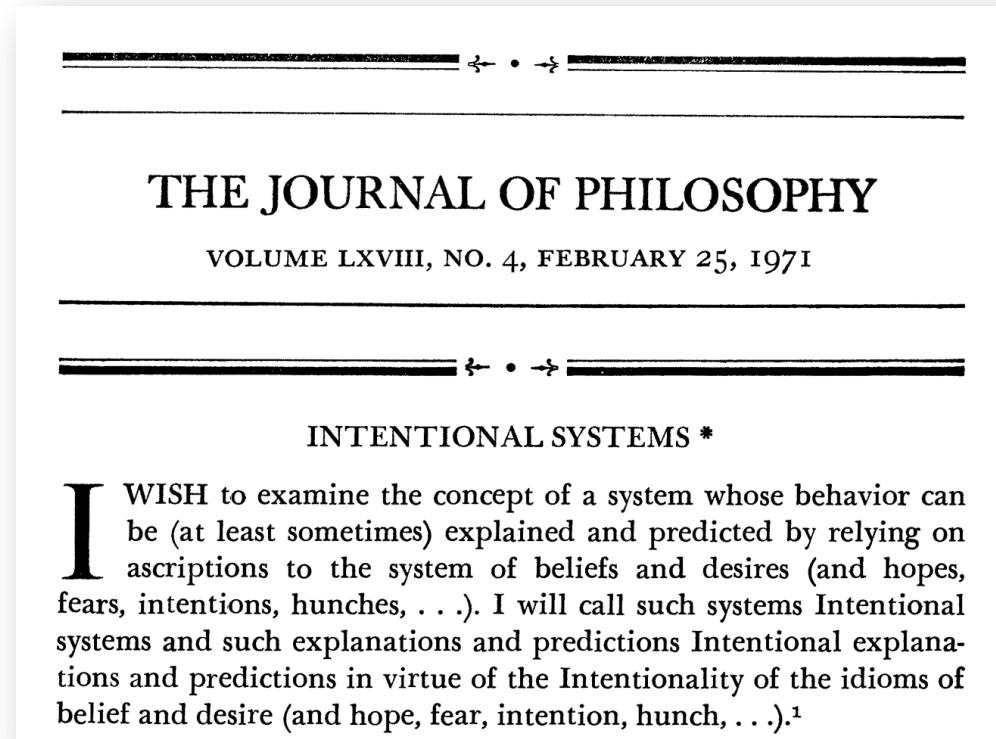
(multi-)agent-oriented => “world of humans”

Agents as **first-class design and programming abstractions**

**Assumption:** ascribing mental qualities to artificial agents simplifies development

Theoretical underpinning:

- Philosophy: the **intentional stance** [Dennett, 1971]



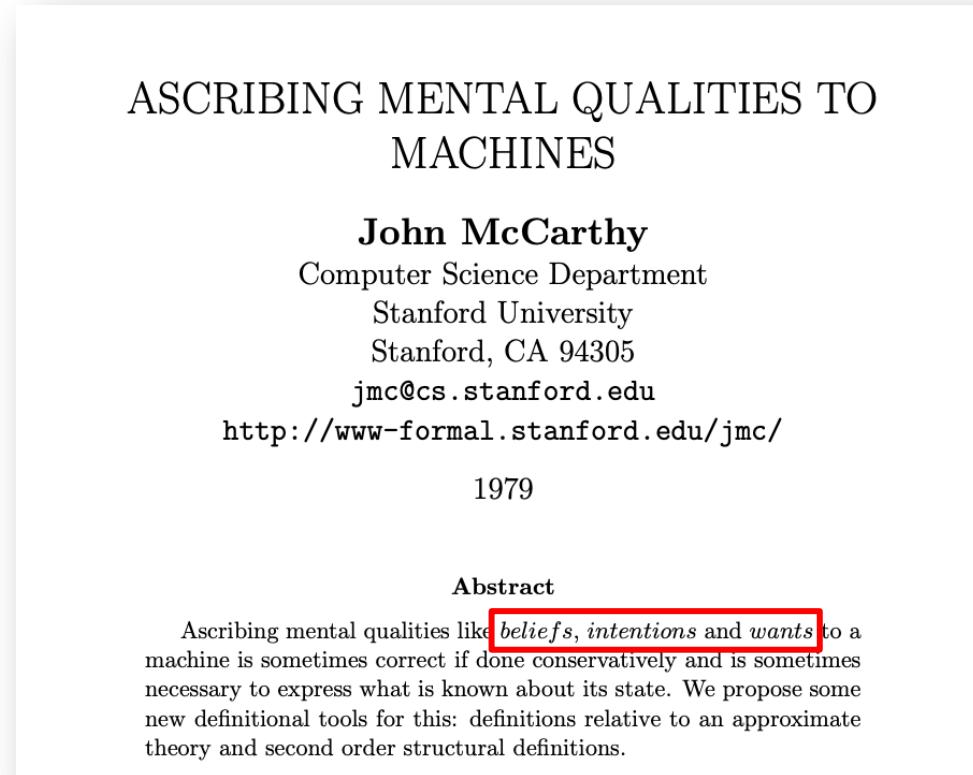
Playing against a **chess program**: do you treat it differently from a human player when trying to **explain or predict** its moves?

Focus is on the **usefulness** of the intentional stance (not if the system has mental qualities legitimately)

Daniel C. Dennett, Intentional Systems, The Journal of Philosophy, Volume 68, Issue 4, 1971.

Theoretical underpinning:

- Philosophy: the **intentional stance** [Dennett, 1971]
- Artificial Intelligence: **ascribing mental qualities** to machines [McCarthy, 1979]



Focus is again on the usefulness of ascribing such mental qualities to machines, not on the legitimacy of the ascription:

“It is useful when the ascription helps us understand the structure of the machine, its past or future behavior, or how to repair or improve it.”

Daniel C. Dennett, Intentional Systems, The Journal of Philosophy, Volume 68, Issue 4, 1971.  
John McCarthy, Ascribing Mental Qualities to Machines, Technical Report ADA071423, Stanford, 1979.

# Belief-Desire-Intention (BDI) Agents

Radical new approach,  
highly influential

## REACTIVE REASONING AND PLANNING

Michael P. Georgeff  
Amy L. Lansky

Artificial Intelligence Center, SRI International  
333 Ravenswood Avenue, Menlo Park, California  
Center for the Study of Language and Information, Stanford University

### Abstract

In this paper, the reasoning and planning capabilities of an au-

This architecture allows PRS to reason about means and ends in much the same way as do traditional planners, but provides the reactivity that is essential for survival in highly dynamic and unstructured environments.

In sum, existing planning systems incorporate many useful techniques for constructing plans of action in a great variety of domains. However, most approaches to embedding these planners in dynamic environments are not robust enough nor sufficiently reactive to be useful in many real-world applications. On the other hand, the more reactive systems developed in robotics are well suited to handling the low-level sensor and effector activities of a robot. Nevertheless, it is not yet clear how these techniques could be used for performing some of the higher-level reasoning desired of complex problem-solving systems. To reconcile these two extremes, it is necessary to develop reactive reasoning and planning systems that can utilize both kinds of capabilities whenever they are needed.

Summary of the related work section  
(the state-of-the-art back in 1987)

Rooted in the BDI model of **human practical reasoning**  
[Bratman, 1987]

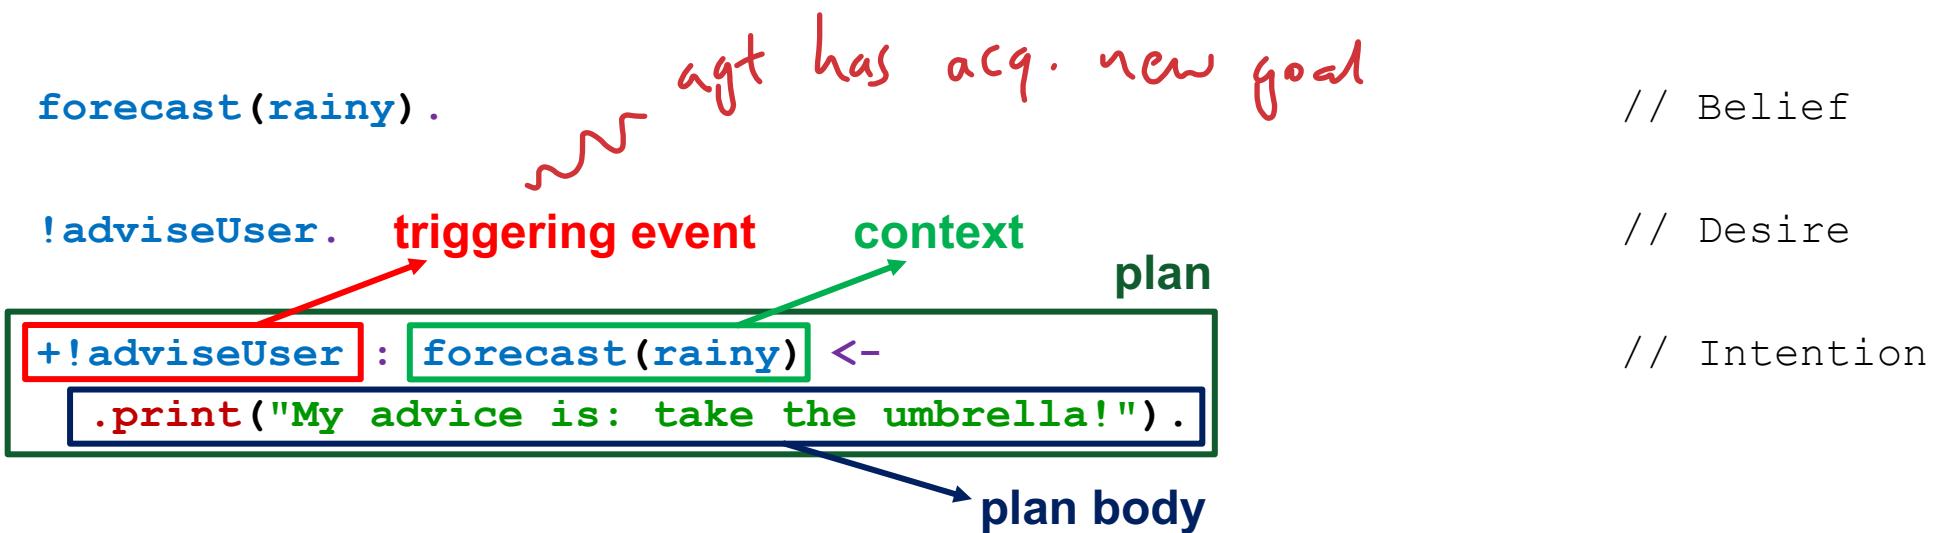
Response to the **diverging views** at the time: logic-based agents vs. reactive agents à la subsumption architecture

**Novelty:** an agent architecture that can balance proactive and reactive behavior (**the Procedural Reasoning System**)

- the internal state of agents is represented in terms of **beliefs**, **desires**, and **intentions** (the agents' **mental states**)
- developers can program agents to **deliberate** about their own beliefs, desires, and intentions — and to modify them as needed:
  - an agent can **decide** to suspend the execution of an intention in order to react to an important event;
  - an agent can **drop an intention** if it becomes unachievable;
  - etc.

M. Georgeff and A. Lansky. *Reactive Reasoning and Planning*. AAAI, 1987.  
Michael Bratman, *Intention, Plans, and Practical Reason*. Harvard University Press, 1987.

# Example (in Jason): Weatherman Bob



- **Beliefs:** information the agent holds about the world; beliefs can be out of date or inaccurate
- **Desires:** states of affairs the agent wishes to bring to the world (broadly, the agent's **goals**)
- **Intentions:** the states of affairs the agent has decided to work towards (goals that the agent is **committed to achieve**)

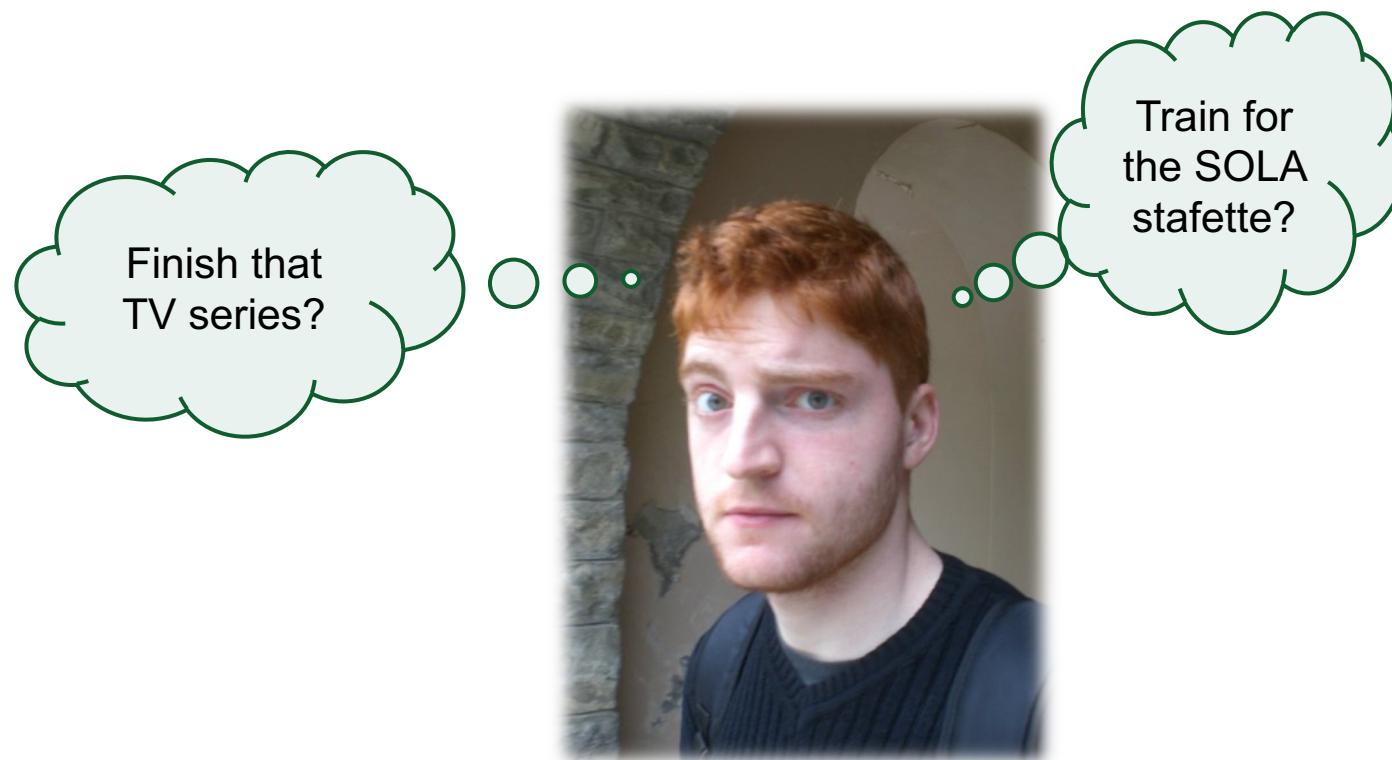
# Today's Agenda

- Introduction
- Reactive Agents
  - The Subsumption Architecture
- Deliberative Agents
  - Agent-Oriented Programming
  - The Belief-Desire-Intention (BDI) Architecture
  - Programming BDI Agents in Jason

# The BDI Architecture

Human practical reasoning involves two main processes [Bratman, 1987]:

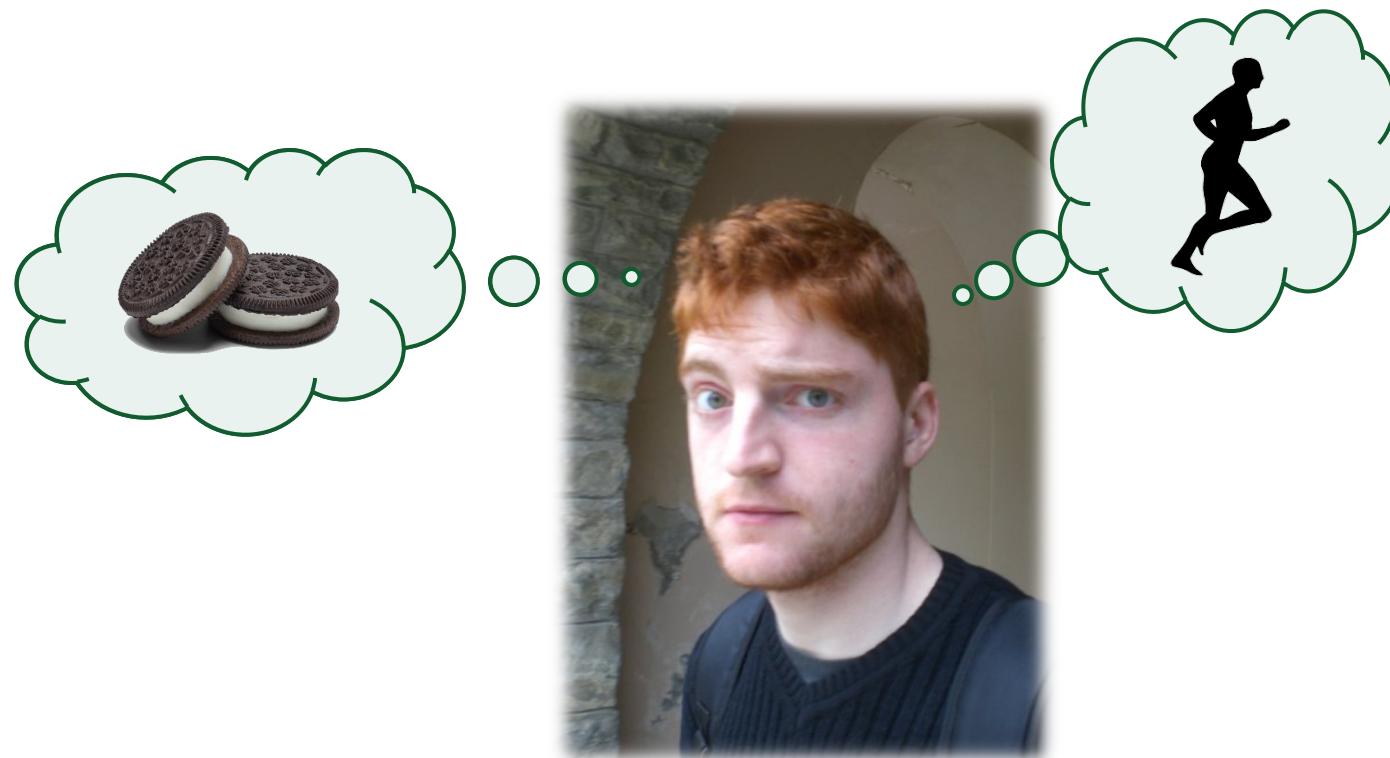
- **deliberation**: deciding *what* goals to achieve
- **means-end reasoning**: deciding *how* to achieve those goals using the available means



Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.  
Michael Bratman, *Intention, Plans, and Practical Reason*. Harvard University Press, 1987.

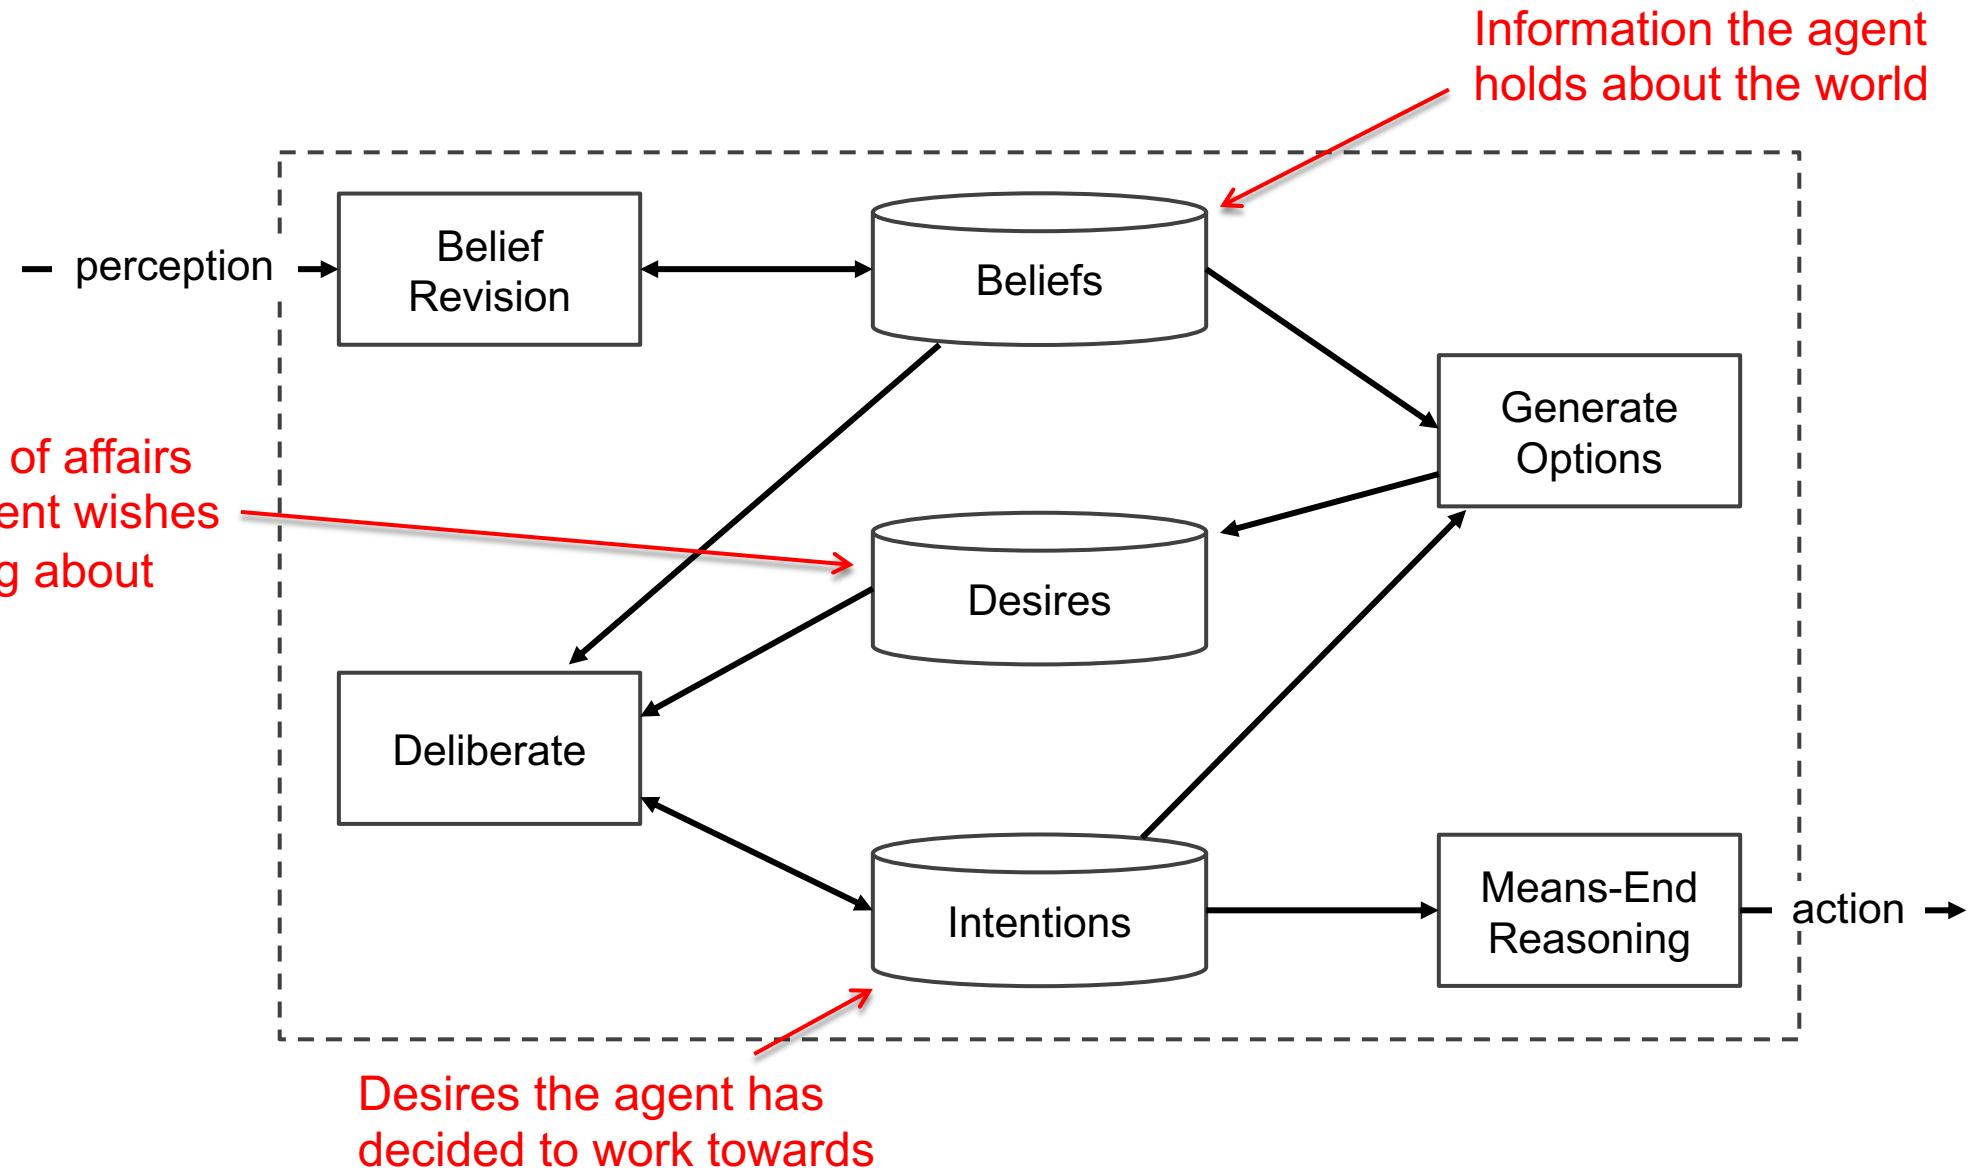
Human practical reasoning involves two main processes [Bratman, 1987]:

- **deliberation**: deciding *what* goals to achieve
- **means-end reasoning**: deciding *how* to achieve those goals using the available means



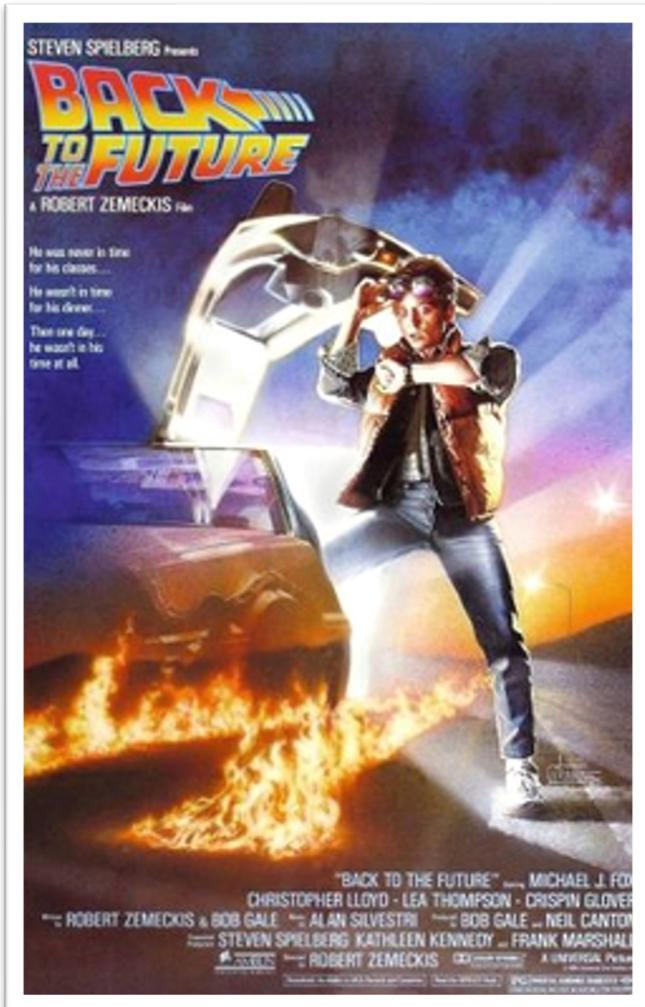
Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.  
Michael Bratman, *Intention, Plans, and Practical Reason*. Harvard University Press, 1987.

# The BDI Architecture: The Mentalistic View



# The BDI Architecture: Reasoning Cycle

```
1 while true do
2   B ← brf(B,perception()) ;                                // belief revision
3   D ← options(B,I) ;                                         // desire revision
4   I ← deliberate(B,D,I) ;                                    // deliberation
5   π ← meansend(B,I,A) ;                                     // means-end reasoning
6   while π != Ø do                                         // (A is the set of
7     execute( head(π) )                                     // available actions)
8   π ← tail(π)
```



“Back to the Future”, 1985  
(from 02:05 to 02:55)

<https://youtu.be/3isQI0nXQRE?t=125>

What's wrong here?

# The BDI Architecture: Reasoning Cycle

```
1 while true do
2   B ← brf(B,perception()) ;                                // belief revision
3   D ← options(B,I) ;                                         // desire revision
4   I ← deliberate(B,D,I) ;                                    // deliberation
5   π ← meansend(B,I,A) ;                                     // means-end reasoning
6   while π != Ø do                                         // (A is the set of
7     execute( head(π) )                                     // available actions)
8   π ← tail(π)
```

Works for proactivity, but not for reactivity (**blind commitment** to the selected plan)

# The BDI Architecture: Reasoning Cycle

```
1 while true do
2   B ← brf(B,perception()) ;                                // belief revision
3   D ← options(B,I) ;                                         // desire revision
4   I ← deliberate(B,D,I) ;                                    // deliberation
5   π ← meansend(B,I,A) ;                                     // means-end reasoning
6   while π != Ø do                                           // (A is the set of
7     execute( head(π) )                                       // available actions)
8     π ← tail(π)
9   B ← brf (B,perception())
10  if ¬sound(π,I,B) then
11    π ← meansend(B,I,A)
```

Revises commitment to the selected plan (re-planning for context adaption)  
... but **blind commitment** to the selected intention  
(never reconsiders the selected intention)

# The BDI Architecture: Reasoning Cycle

```
1 while true do
2   B ← brf(B,perception()) ;                                // belief revision
3   D ← options(B,I) ;                                         // desire revision
4   I ← deliberate(B,D,I) ;                                    // deliberation
5   π ← meansend(B,I,A) ;                                     // means-end reasoning
6   while π != Ø and ¬succeeded(I,B) and ¬impossible(I,B) do // (A is the set of
7     execute( head(π) )                                         // available actions)
8   π ← tail(π)
9   B ← brf (B,perception())
10  if ¬sound(π,I,B) then
11    π ← meansend(B,I,A)
```

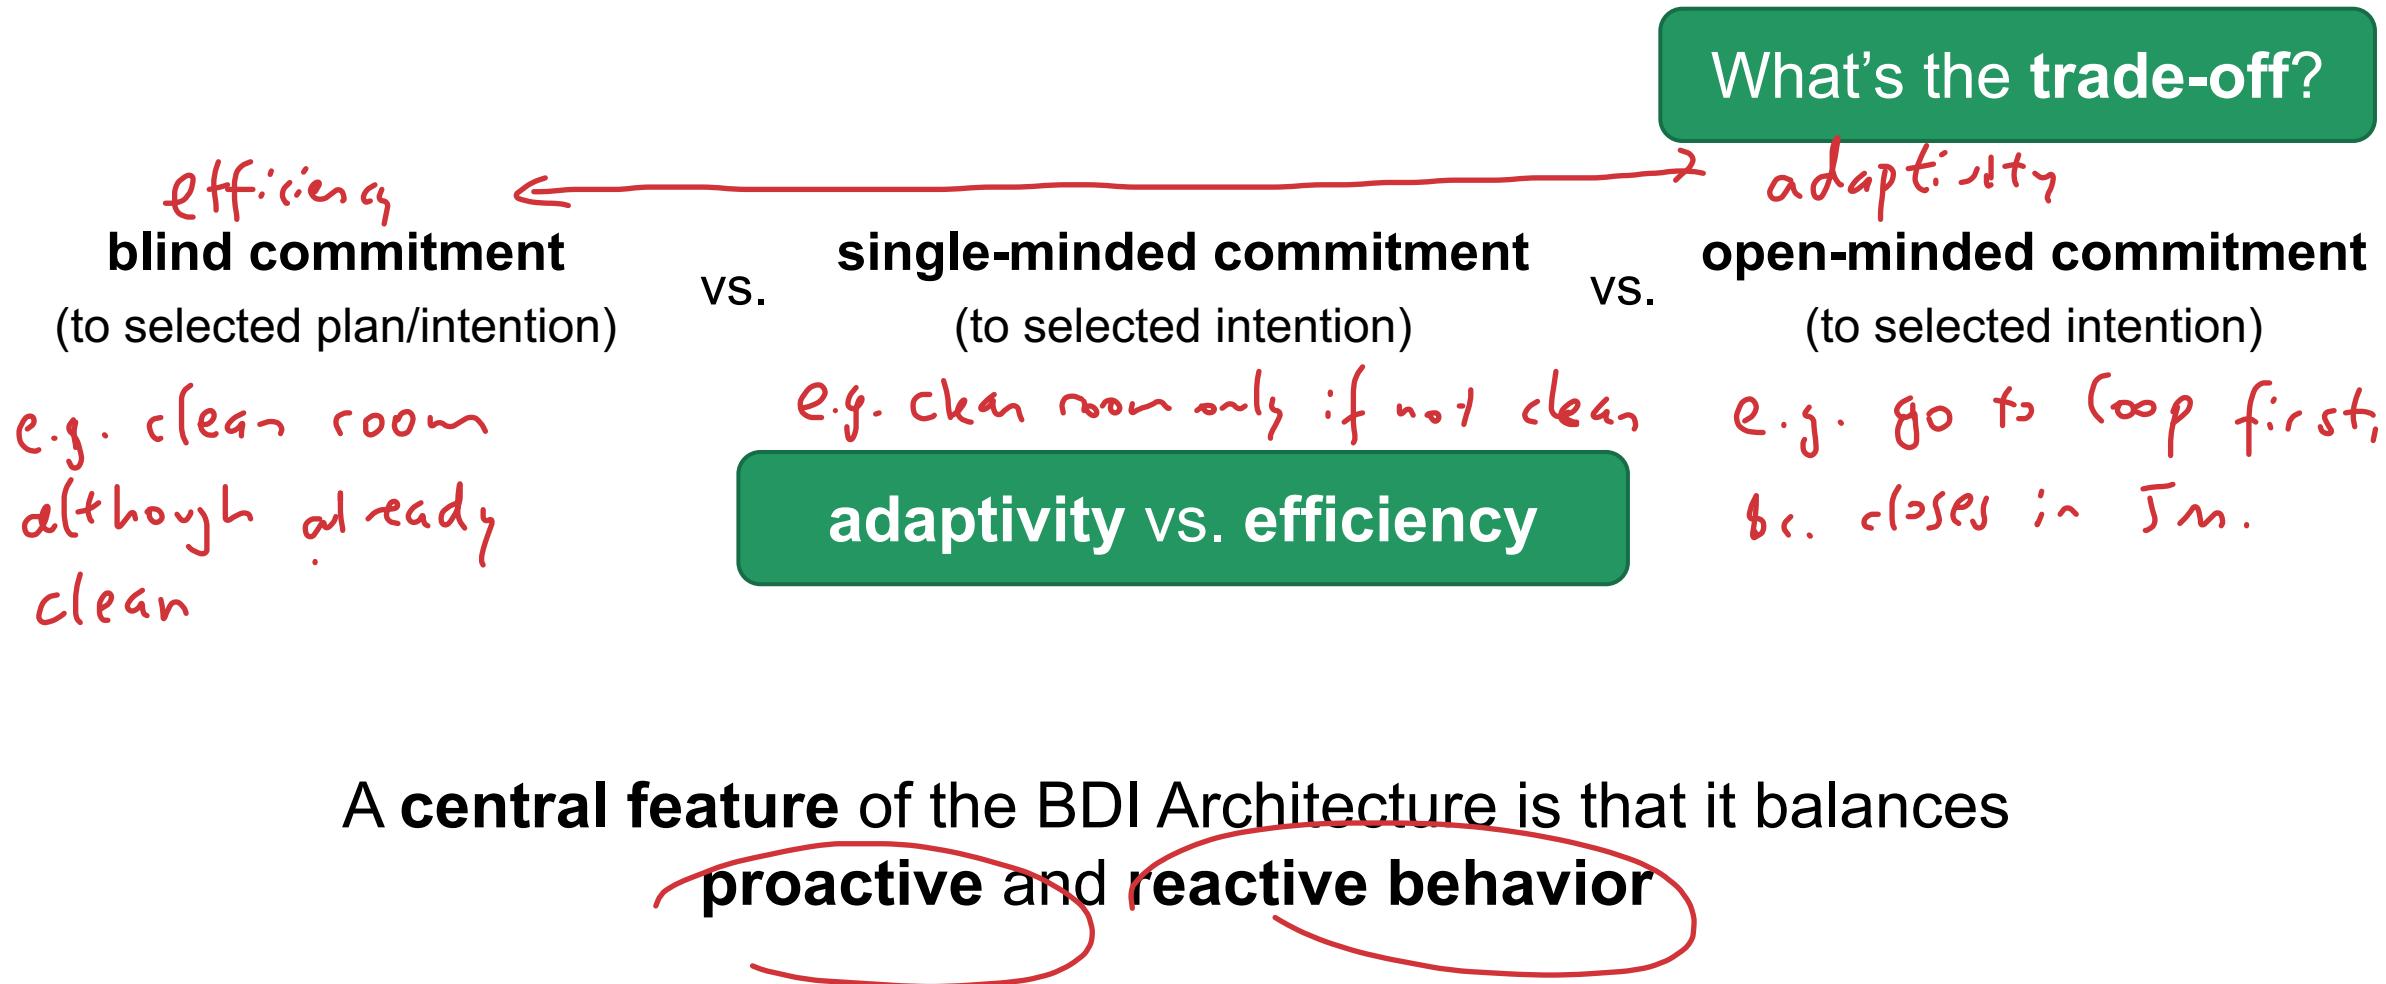
Checks if the selected intention was already achieved or became impossible  
... but **single-minded commitment** to the selected intention  
(will not drop the selected intention while it is achievable)

# The BDI Architecture: Reasoning Cycle

```
1 while true do
2   B ← brf(B,perception()) ;                                // belief revision
3   D ← options(B,I) ;                                         // desire revision
4   I ← deliberate(B,D,I) ;                                    // deliberation
5   π ← meansend(B,I,A) ;                                     // means-end reasoning
6   while π != Ø and ¬succeeded(I,B) and ¬impossible(I,B) do // (A is the set of
7     execute( head(π) )                                         // available actions)
8   π ← tail(π)
9   B ← brf (B,perception())
10  if reconsider(I,B) then
11    D ← options(B,I)
12    I ← deliberate(B,D,I)
13  if ¬sound(π,I,B) then
14    π ← meansend(B,I,A)
```

An **open-minded** BDI agent can reconsider its intentions (not necessarily at each iteration)  
(only maintains its intentions if they are still its desires)

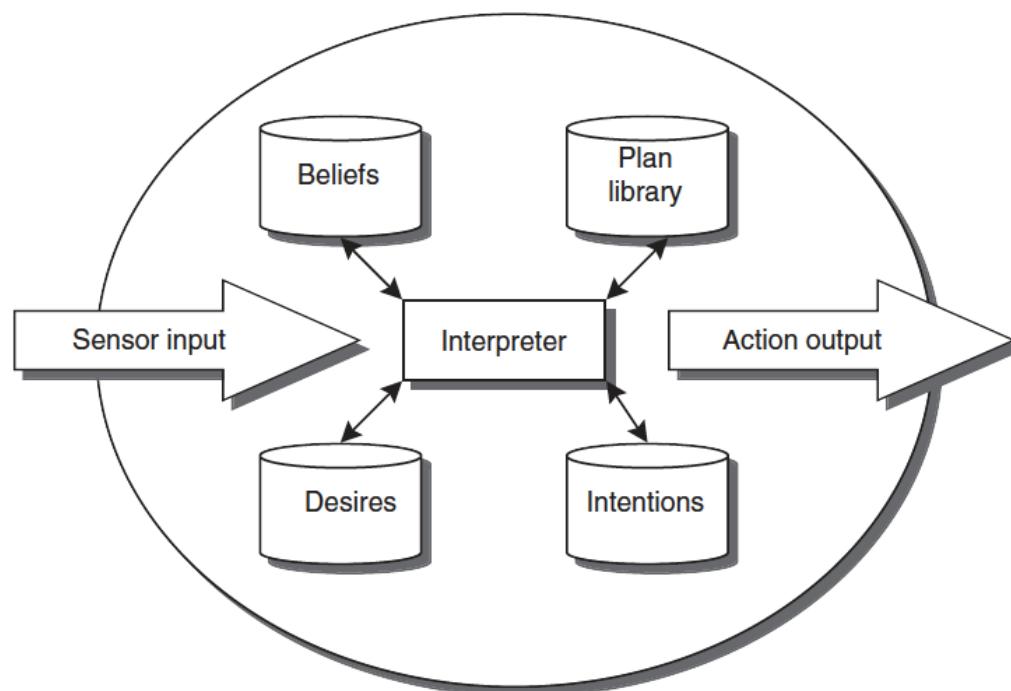
# The BDI Architecture: Reasoning Cycle



# The BDI Architecture: Means-End Reasoning

**Means-end reasoning:** the process of deciding how to achieve an end (i.e., an intention) using the available means (i.e., the actions that can be performed in the environment)

- **automated planning:** yields proactive behavior, but low reactivity
- **Procedural Reasoning System (PRS)** [Georgeff & Lansky, 1987]: balances proactivity and reactivity



Radical new idea: **procedural reasoning**

- no automated planning
- agents are equipped with libraries of plans (**procedural knowledge**)
- agents assemble plans at run time (**procedural reasoning**)

To program such agents, we need a  
**practical programming language**

# Today's Agenda

- Introduction
- Reactive Agents
  - The Subsumption Architecture
- Deliberative Agents
  - Agent-Oriented Programming
  - The Belief-Desire-Intention (BDI) Architecture
  - Programming BDI Agents in Jason

# Programming BDI Agents in Jason

**Jason** is an interpreter for an extended version of AgentSpeak(L)

- AgentSpeak(L) [Rao, 1996] is a programming language for BDI agents inspired by the PRS and grounded in first-order logic
- Jason brings several practical extensions such as communication

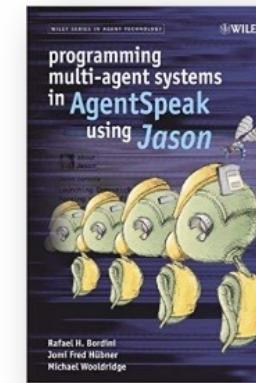
Jason is part of the JaCaMo platform [Boissier et al., 2020]

- Jason commonly used to refer to the programming language it implements

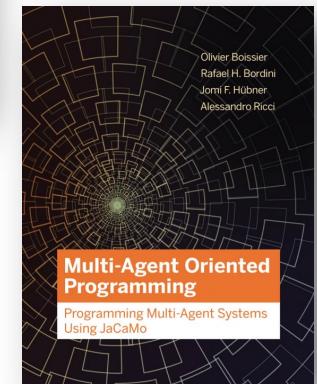
Main language constructs: beliefs, goals, and plans



Abstractions so intuitive they could work for non-technical users



Chapters 1-3



Chapters 3-4

Anand S. Rao, *AgentSpeak(L): BDI agents speak out in a logical computable language*, MAAMAW 1996.  
 Rafael Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.  
 O. Boissier, R. H. Bordini, J.F. Hubner, A. Ricci. *Multi-Agent Oriented Programming: Programming Multi-Agent Systems Using JaCaMo*, The MIT Press, 2020.  
 Samuele Burratini et al., *Agent-Oriented Visual Programming for the Web of Things*, EMAS 2022 (in submission).

Beliefs are represented by annotated predicates of first-order logic:

```
predicate(term1,...,term2) [annot1,...,annot2]
```

Beliefs can be inaccurate, outdated, contradictory,  
may have different degrees of certainty, etc.

## Example: Weatherman bob's belief base

```
forecast(rainy) [source(percept),certainty(0.6)] .  
forecast(sunny) [source(jane),certainty(0.8)] .
```

// beliefs about the environment

```
knows(bob, jane) [source(self)] .  
knows(bob, joe) [source(self)] .
```

// beliefs about the agent itself

```
sincere(jane) [source(self)] .  
~sincere(joe) [source(self),source(jane)] .
```

// beliefs about other agents

Strong negation is used to represent what is believed to  
be false (necessary under the open-world assumption)

Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.

Jomi F. Hübner, Multi-Agent Oriented Programming Using JaCaMo, AI4Industry Summer School, MINES Saint-Étienne, 2020.

A Jason agent starts with an initial set of beliefs, i.e. its **initial belief base** (can be empty)

- the agent can acquire new beliefs throughout its lifetime

## Through perception

An agent can perceive its environment to acquire new beliefs (annotated with `source(percept)`)

- beliefs based on perception are updated automatically during reasoning cycles by the Jason interpreter

```
forecast(rainy) [source(percept),certainty(0.6)].
```

## Through intention

An agent can create/delete beliefs (**mental notes**) as part of a plan using the + and – operators (annotated with `source(self)`)

```
+sincere(jane);           // adds sincere(jane) [source(self)] to the belief base
-sincere(jane);           // removes sincere(jane) [source(self)] from the belief base
```

A Jason agent starts with an initial set of beliefs, i.e. its **initial belief base** (can be empty)

- the agent can acquire new beliefs throughout its lifetime

More on **agent communication languages** next week!

## Through communication

When an agent receives a **tell** message, the content is added to the belief base as a new belief annotated with the sender of the message as the source

```
.send(bob, tell, forecast(rainy) [certainty(0.6)]); // sent by jane  
// adds forecast(rainy) [source(jane),certainty(0.6)] in bob's belief base
```

If an agent receives an **untell** message and the content of the message matches a belief in the agent's belief base with the sender as the source, then the belief is removed

```
.send(bob, untell, forecast(rainy) [certainty(0.6)]); // sent by jane  
// removes forecast(rainy) [source(jane),certainty(0.6)] from bob's belief base
```

Rules can be used to infer new beliefs:

```
happy(bob) :- today(saturday) | today(sunday) .
```

```
advice("take the umbrella") :- forecast(rainy)[source(Ag),certainty(Cert)] &
sincere(Ag) & Cert >= 0.1 .
```

## Useful Operators

- boolean operators: & (and), | (or), not (not), = (unification), relational operators (>, >=, <, <=), == (equal), \== (different)
- arithmetic operators: + (sum), – (subtraction), \* (multiply), / (divide), div (division quotient), mod (division remainder), \*\* (power)

Types of goals:

- achievement goals: goals **to do**

```
!giveWeatherAdvice.           // bob's initial goal
```

- test goals: goals **to know**; can be used inside plan bodies to query the belief base

```
?advice(Advice) ;           // test goal for retrieving an advice from the belief base
```

Same syntax as beliefs, but goals are preceded by  
! (achievement goal) or ? (test goal)

# Goals: Dynamics

A Jason agent starts with an **initial set of goals** (can be empty), and can acquire new goals throughout its lifetime

## Through intention

The ! and ? operators can be used inside a plan to create a new goal with `source(self)`

```
// creates a new achievement goal !giveWeatherAdvivce[source(self)]  
!giveWeatherAdvivce;  
// creates a new test goal ?advice(Advice) [source(self)]  
?advice(Advice);
```

Both achievement goals and test goals can also be acquired through communication

More on **agent communication languages** next week!

Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.

Jomi F. Hübner, Multi-Agent Oriented Programming Using JaCaMo, AI4Industry Summer School, MINES Saint-Étienne, 2020.

A Jason plan has the following structure:

```
triggering_event : context <- body.
```

where:

- the **triggering event** denotes the event that the plan is meant to handle
- the **context** represents the circumstances in which the plan can be used (**applicability precondition**)
- the plan's **body** is the course of action to be used for handling the event if the context is believed to be true at the time a plan is chosen to handle the event



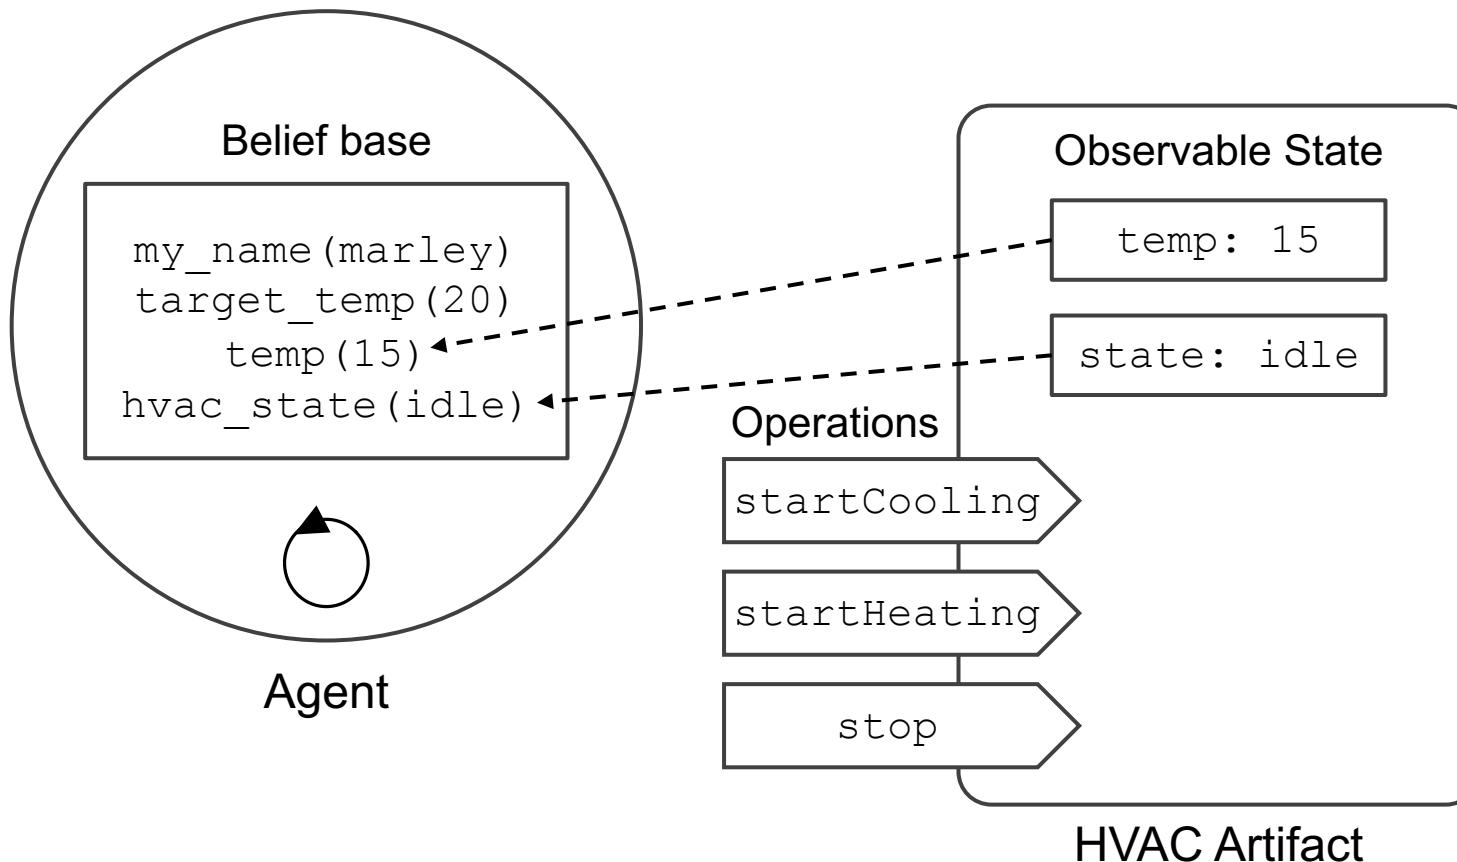
## Triggering Events

Events are created as consequences to changes in an agent's beliefs or goals:

- `+b` (belief addition)
- `-b` (belief deletion)
- `+ !g` (achievement-goal addition)
- `- !g` (achievement-goal deletion)
- `+ ?g` (test-goal addition)
- `- ?g` (test-goal deletion)

An agent reacts to events by executing plans

# Triggering Events: Proactive vs. Reactive Plans



**Proactive plans:** pursue (achievement) goals

**Reactive plans:** react to environmental changes

## Proactive Plans

```
+!set_temp(Target) : temp(Current)
  & Target > Current
  <- startHeating.

+!set_temp(Target) : temp(Current)
  & Target < Current
  <- startCooling.

+!set_temp(Target) : temp(Current)
  & Target == Current
  <- stop.
```

## Reactive Plans

```
// maintenance task
+temp(Current)
  : target_temp(Target)
  & Target \== Current
  <- !set_temp(Target).
```

Useful operators for the plan's context (same as for inference rules):

- boolean operators: & (and), | (or), not (not), = (unification), relational operators (>, >=, <, <=), == (equal), \== (different)
- arithmetic operators: + (sum), – (subtraction), \* (multiply), / (divide), div (division quotient), mod (division remainder), \*\* (power)

Useful operators for the plan's body:

```
+rain : time_to_leave(T) & time_now(H) & H <= T <-
    !new_goal;                                // new sub-goal
    ?distance_to_destination(X);              // new test goal
    X > 10;                                    // constraint to carry on
    X = 100;                                   // assignment
    +note(T - H);                            // add mental note
    -note(T - H);                            // remove mental note
    -+note(T * H);                           // update mental note
    .print("Hello world! Time to leave is: ", T); // internal action (a print statement)
    open(door).                                // external action
```

Reference to internal actions provided by Jason: <http://jason.sourceforge.net/api/jason/stdlib/package-summary.html>

A Jason agent is equipped with a **library of plans**

The plan library is initialized with the plans defined by the programmer of the agent (in the agent program), but can **evolve** at run time:

- **through intention:** plans can be added or removed dynamically and intentionally by the agent using two internal actions: `.add_plan` and `.remove_plan` (see reference in footnote)
- **through communication:** agents can exchange plans just like they can exchange beliefs or delegate goals

More on **agent communication languages** next week!

Jason agents could also synthesize plans, but this functionality is not trivial — and is not provided out-of-the-box by the JaCaMo platform

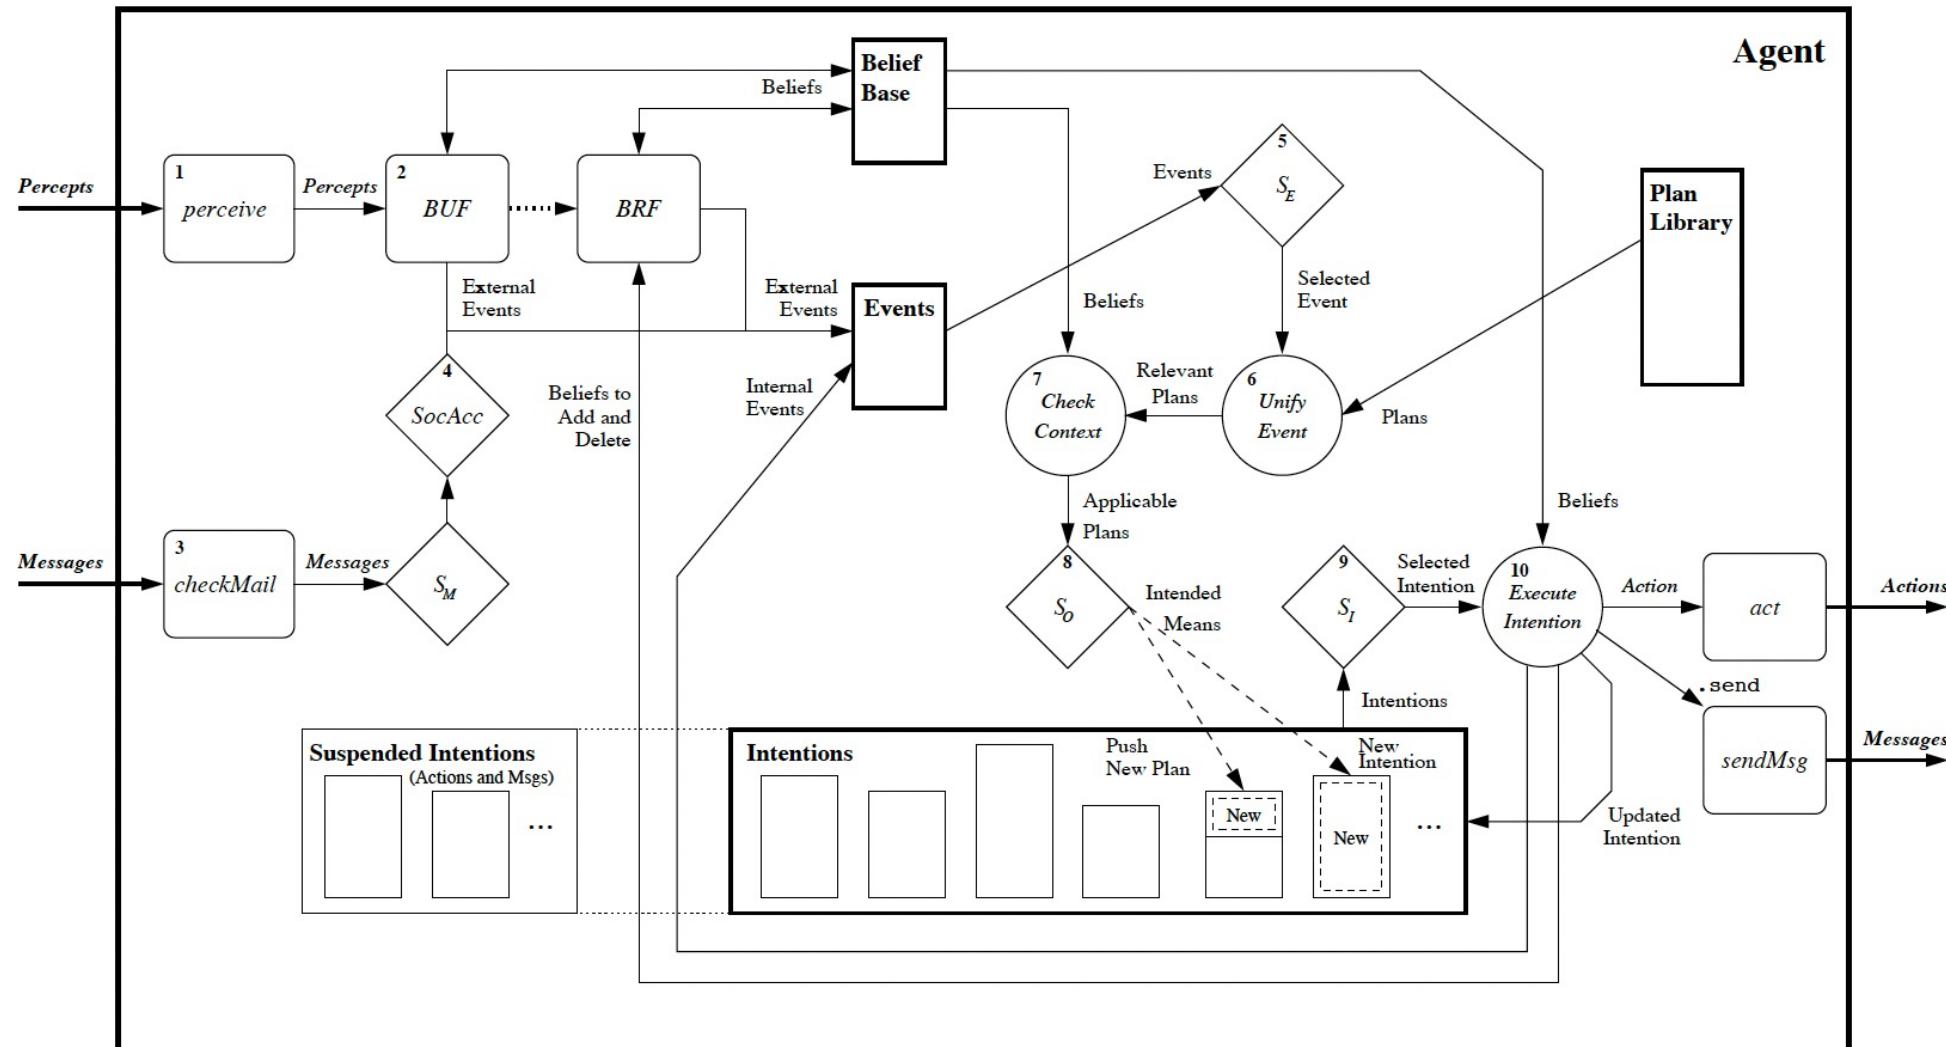
- ex: agents could use an external service for PDDL-based planning

Reference to the internal actions supported by Jason: <http://jason.sourceforge.net/api/jason/stdlib/package-summary.html>

Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.

Jomi F. Hübner, Multi-Agent Oriented Programming Using JaCaMo, AI4Industry Summer School, MINES Saint-Étienne, 2020.

# A Glimpse of the Jason Reasoning Cycle

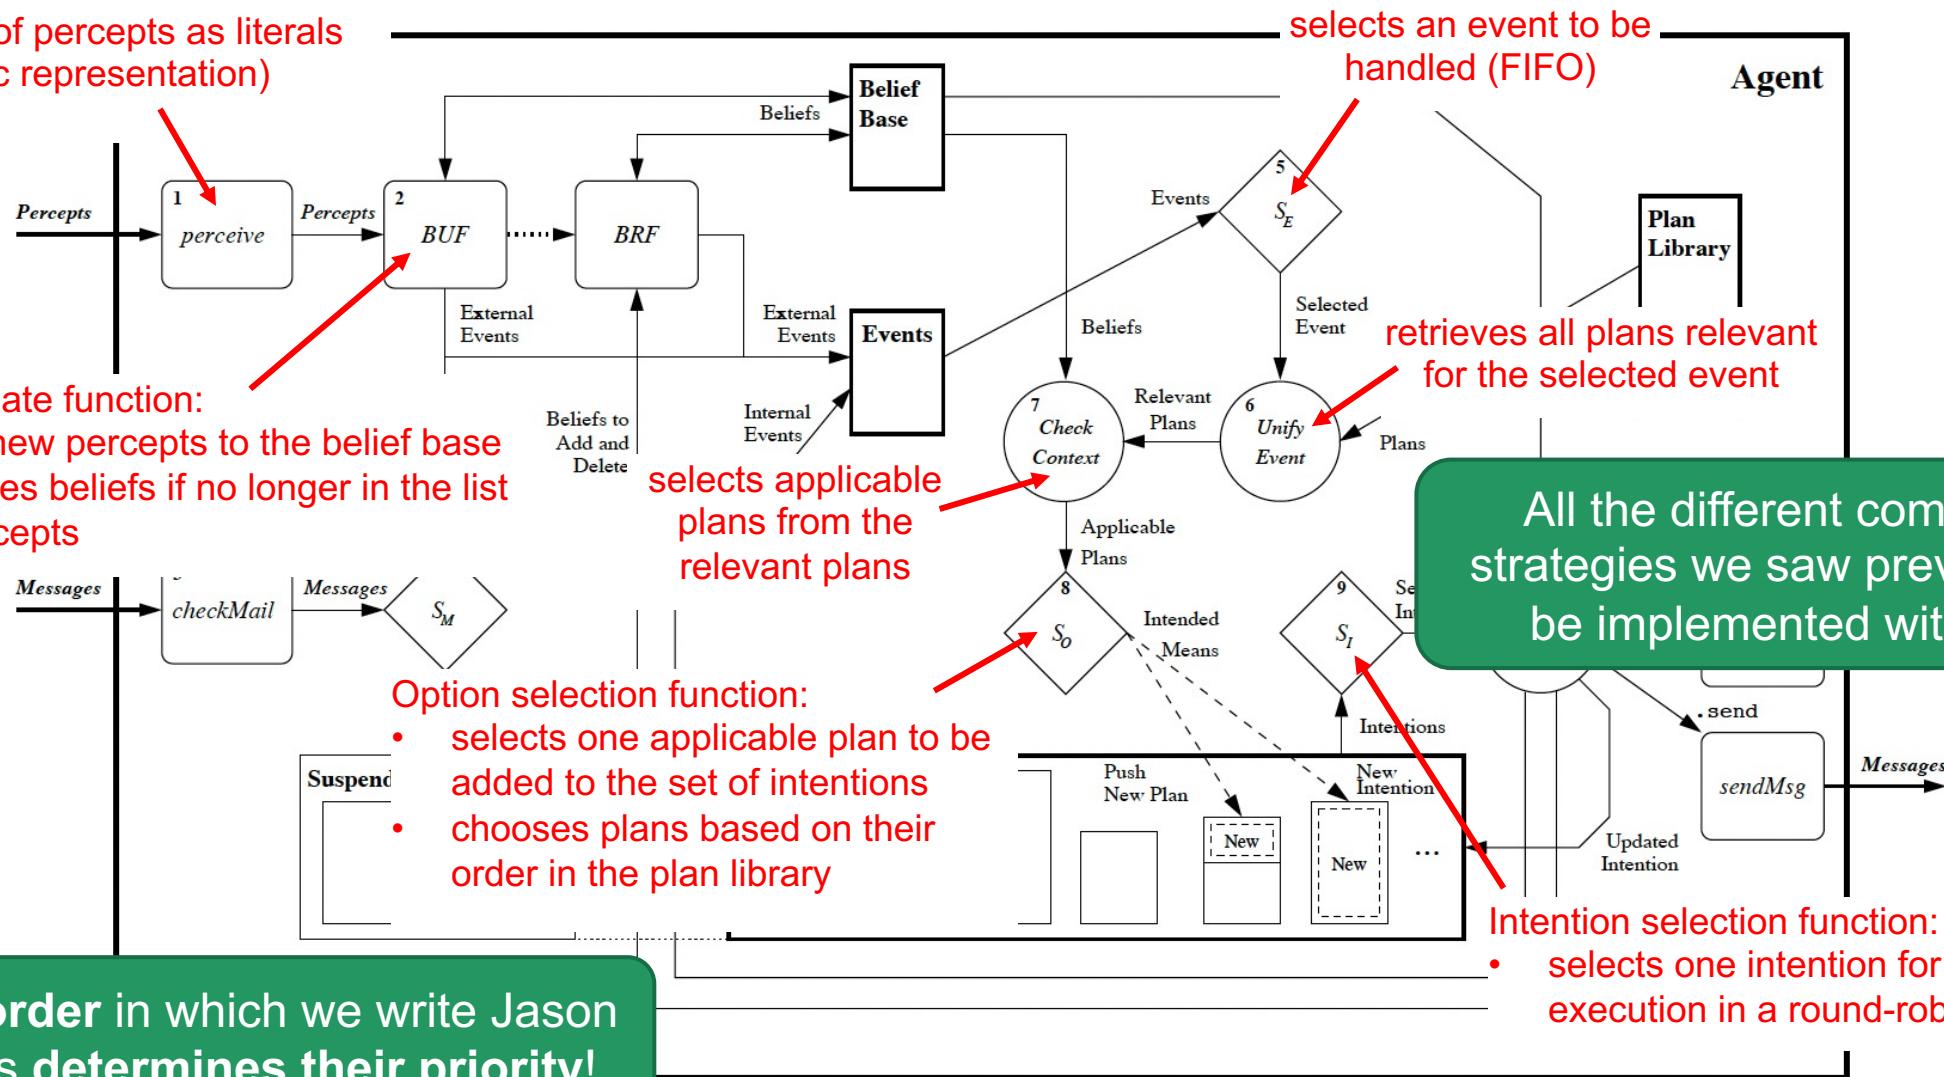


Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.

Jomi F. Hübner, Multi-Agent Oriented Programming Using JaCaMo, AI4Industry Summer School, MINES Saint-Étienne, 2020.

# A Glimpse of the Jason Reasoning Cycle

returns a list of percepts as literals  
(symbolic representation)



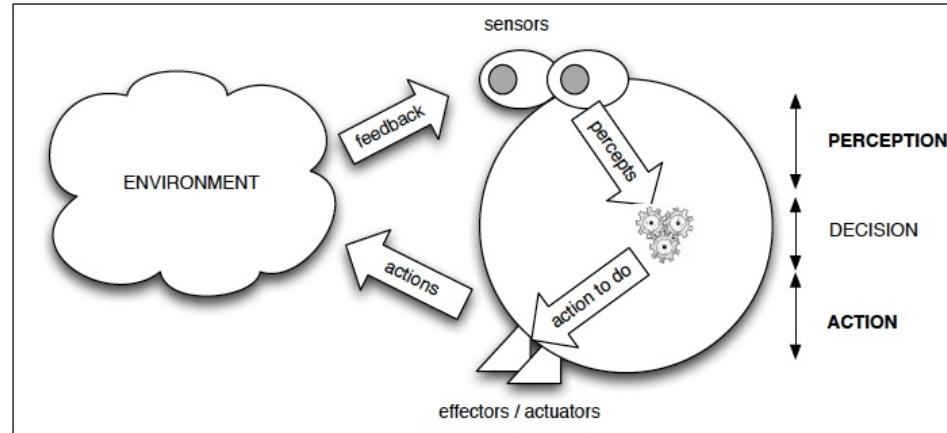
The **order** in which we write Jason plans determines their priority!

Rafel Bordini et al., *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.

Jomi F. Hübner, Multi-Agent Oriented Programming Using JaCaMo, AI4Industry Summer School, MINES Saint-Étienne, 2020.

# Concluding Remark: Agents vs. Objects/Actors

Agents as first-class design/programming abstractions:



- task/goal-oriented
- pro-active + reactive
- decision making

## Agents vs. Objects

- agents are active, objects are passive (except for active objects)
- agents have stronger encapsulation:
  - state + behavior **+ control of their behavior**
  - decision making

## Agents vs. Actors

- actors are reactive, agents are also proactive
- actors operate on an event-loop, agents operate on a reasoning cycle
- actors are message-driven, agents are task/goal-directed

# Our Journey

**Prerequisites:**  
 • ASSE  
 • (...)



**Week 1:  
Introduction**



**Week 2:  
A Web for Machines**



**Week 3:  
Knowledge Representation  
and Reasoning for the Web**



**Week 4:  
Linked Data and Distributed  
Knowledge Graphs**



**Week 5:  
Autonomous Agents  
(Arch. and Programming)**



**Week 10:  
Game Theory and  
Social Choice**



**Week 6 (Coordination I):  
Agent Communication  
and Interaction**



**Week 9 (Coordination IV):  
Trust & Reputation**



**Week 11:  
Reinforcement Learning  
and Multi-Agent Learning**



**Week 12:  
An Industry Perspective**



**Exercises**

Ex1: Writing Your First Agent(s)!  
 Ex2: Automated Planning  
 Ex3: Web Ontologies  
 Ex4: Operating on Linked Data  
 Ex5: BDI Agents  
 Ex6: Interacting Agents on the Web

Ex7: Ant Colony Optimization  
 Ex8: Organized Agent  
 Ex9: Trustworthy Agents  
 Ex10: Axelrod's Agents  
 Ex11: Reinforcement Learning Agents  
 Course Review and Q&A

# Any Questions / Comments / Doubts / Concerns?



<https://www.istockphoto.com/>

<https://freepik.com>